



# Deep Learning - MAI



Convolutional neural networks

THEORY

Dario Garcia Gasulla  
*dario.garcia@bsc.es*

# Spatial Connectivity

- ❖ Some data has spatial correlations that can be exploited
  - 1D, 2D, 3D, ...
- ❖ Near-by data points are more relevant than far-away.
- ❖ Sparsify connectivity to reduce complexity and ease the learning



# Weight Sharing

- ❖ Sparse connectivity is nice, but we want to apply filters everywhere.
- ❖ Each filter will get convolved all over the image: 2D activations matrix
- ❖ In static we have sets of neurons sharing weights
- ❖ In this context, what is a neuron?



# Convolution in Action

Kernel size 3x3  
(neuron input = 9)

1 0 1  
0 1 0  
1 0 1

Detect 'X'

1 <sub>x1</sub>	1 <sub>x0</sub>	1 <sub>x1</sub>	0	0
0 <sub>x0</sub>	1 <sub>x1</sub>	1 <sub>x0</sub>	1	0
0 <sub>x1</sub>	0 <sub>x0</sub>	1 <sub>x1</sub>	1	1
0	0	1	1	0
0	1	1	0	0

Image

Filter convolution  
process

4		

Convolved  
Feature

Activations (pre-func.)

# Image Transformations

- ❖ Convolving filters transform the image
- ❖ Let the model learn the kernels it needs

**Edge detection**

$$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$$




**Sharpen**

$$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$$



**Gaussian blur**  
 $3 \times 3$

$$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$$




# Convolution Details

**Kernel size:** Size of the receptive field of convolutional neurons

**Stride:** Steps size of convolution

**Padding:** Allows focus on border

- ❖ Most common fill: Zeros
- ❖ Valid (no padding): Internal only. May skip data. Reduces dimensionality
- ❖ Same: Keep dimensionality with stride 1




$$OutputSize = \frac{InputSize - KernelSize + 2 * Padding}{Stride} + 1$$

# Dilated/Atrous Convolutions

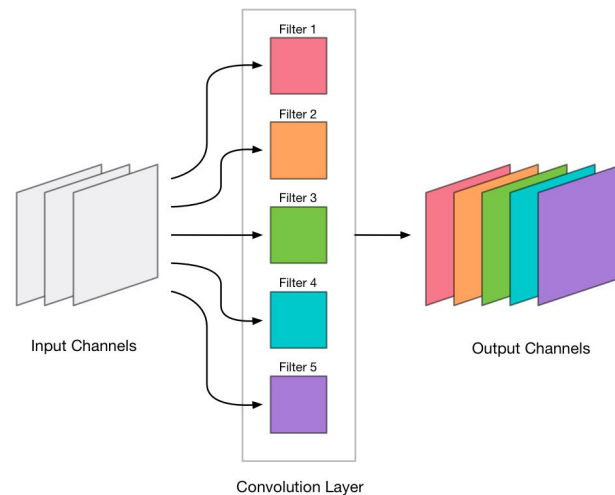

Sparsify the kernel

- ❖ Increases perceptive field without added complexity
- ❖ Loses details, gains context
- ❖ Another hyperparam :(
- ❖ Used for
  - Down/Upsampling (segmentation)
  - High Resolution inputs



# Output Volumes

- ❖ Typically, conv filters are full depth ( $N*N*input\_depth$ )
- ❖ Each conv filter (often 3D) convolved generates a 2D plane of data
- ❖ Depth provides all the *views* on a part of the input
- ❖ Output volume: New representation of input with different dimensions



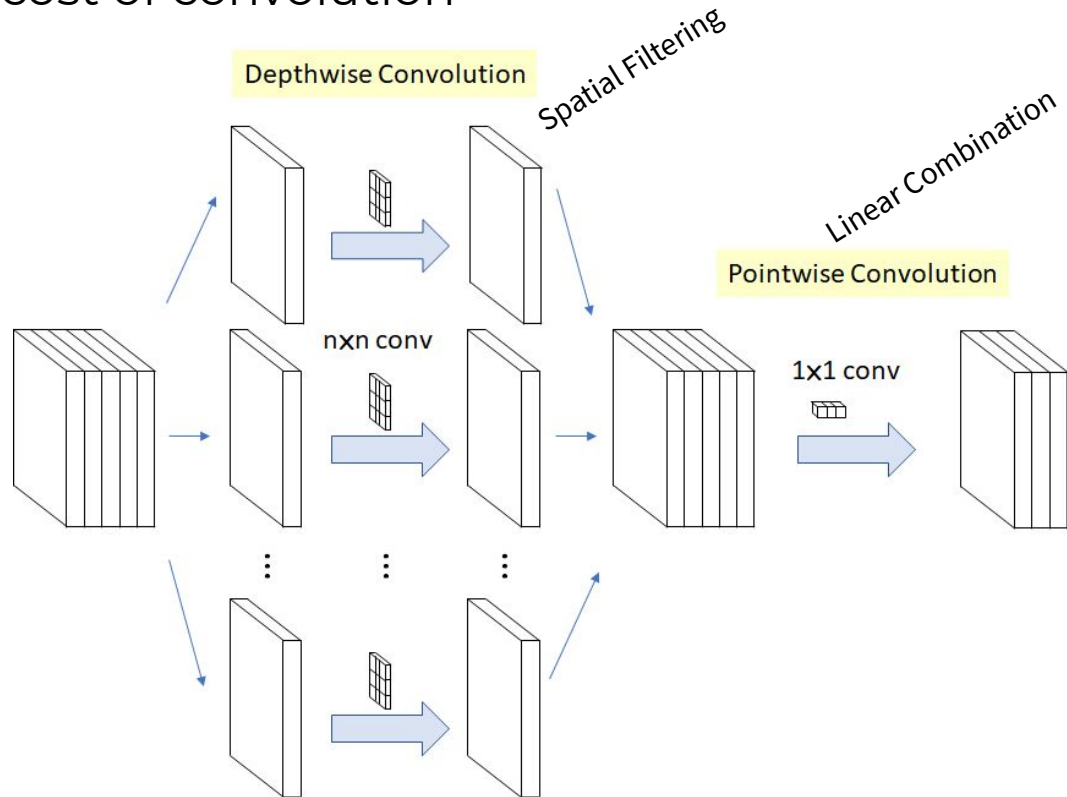


# Depth-wise Separable Convolutions

Decreasing the complexity and cost of convolution

1. Depth-wise convolutions
  - Filters:  $N \times N \times 1$
2. Point-wise convolution
  - Filters:  $1 \times 1 \times \text{input\_depth}$

Params:  $N \times N + N$





# To Pool Or Not To Pool

- ❖ Operation: **Max** or Avg
- ❖ Dimensionality reduction (along x and y only)
- ❖ Rarely applied full depth
- ❖ Parameter free layer
- ❖ Hyperparams: Size & Stride
- ❖ Loss in spatial precision / Robust to invariance


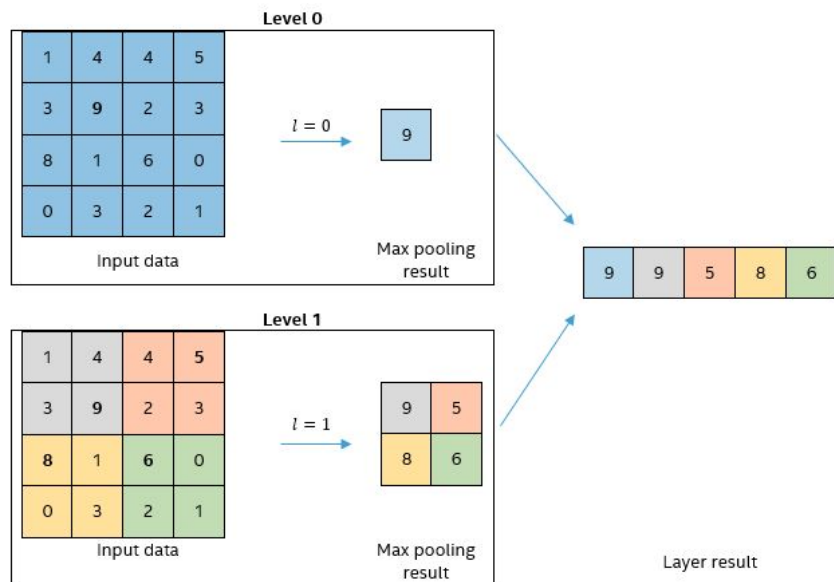
Other means to reduce complexity

- ❖ Depth-wise separable convs, bigger conv. strides



# Spatial Pyramid Pooling (SPP)

- ❖ Multi-scale Pool (by powers of 2)
- ❖ Often used between conv and fc



More alternatives: Atrous spatial PP, Global average pooling, Pyramid pooling module, Adaptive PP

# Practical Tips XI

## Convolutional

- Small/big filters (**3x3**, 5x5, 7x7)
  - Cheap/Expensive
  - Local/General
  - Bigger/Smaller outputs (stride)
- Kernel Size = input size: fc
- Kernel size = 1x1: Alter depth)

## Pooling

- 2x2, stride 1 is the least invasive


## Hyperparameters incomplete list #4

- ☐ Kernel size (conv & pool)
- ☐ Stride (conv & pool)
- ☐ Padding (conv & pool)
- ☐ Num. filters
- ☐ Dilatation rate

# The Challenge

## ImageNet Large-Scale Visual Recognition Challenge 2012 (ILSVRC'12)

- ❖ Image Classification: 1,000 classes
- ❖ Training: 1.2M
- ❖ Living things + Human-made objects
  - 120 breeds of dogs



# Data Augmentation for CNNs

Apply what is safe for each case



- ❖ Horizontal/Vertical flips
- ❖ Random Crops
- ❖ Color Jitter
- ❖ Rotation
- ❖ ...



# CNNs Big Bang

## AlexNet (2012)

- ❖ Breakthrough in ILSVRC
- ❖ 5 convs+pools, ReLU, 2 dense, and dropout
- ❖ 62M parameters





On the shoulders of giants

# A new Standard for CNNs

VGG 11/13/16/19 (2014)

- ❖ Prototype of (conv-pool)\*+dense\* architecture
- ❖ 133-144M parameters
- ❖ 3x3 convs only

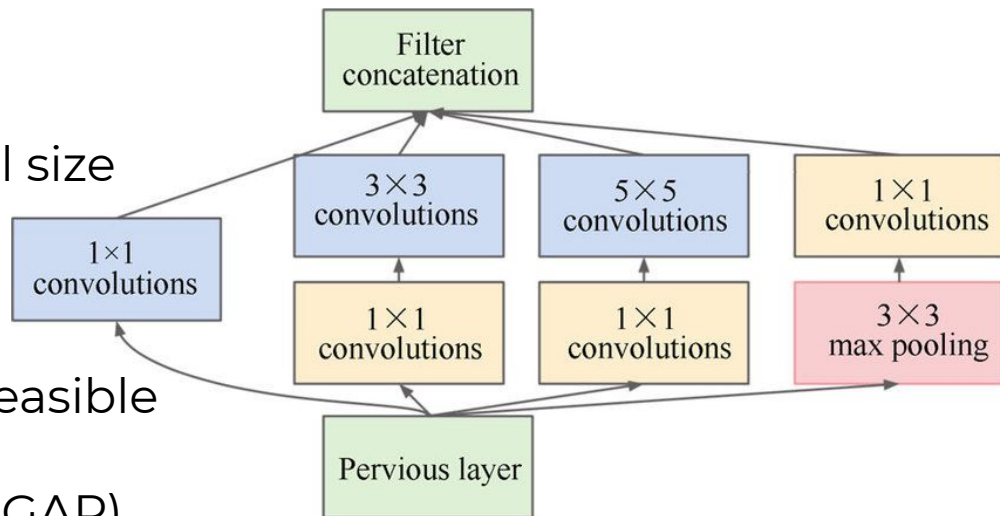




# The Inception Family

GoogLeNet (2014)

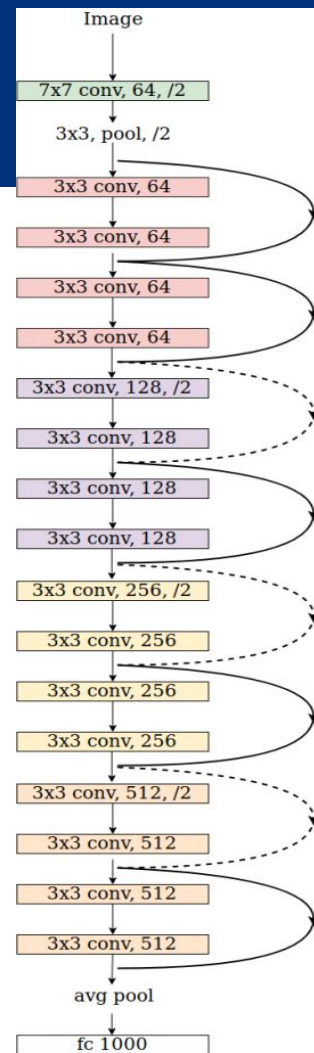

- ❖ The Inception block
- ❖ Let the model decide the kernel size
- ❖ Better scale adaptation
- ❖ Bottleneck  $1 \times 1$  conv to make it feasible
- ❖ No FC: Global Average Pooling (GAP)



# The Skipped Connection

## ResNet (2015)

- ❖ Residual blocks / Skip connections
- ❖ Deeper should never be worse
  - Learning the identity is hard
  - Learning to cancel out is easy
- ❖ Shallow ensemble of nets
- ❖ Train up to 1K layers (do not!)
- ❖ ILSVRC'12 human level




# Inverted Residuals and Linear Bottlenecks


1. Upsample depth
2. Depth-wise conv
3. Point-wise conv

- ❖ Linear act at end
- ❖ Non-linear mid
- ❖ Residual link
- ❖ Efficient

(a) Residual block



(b) Inverted residual block



*Sponsored by:*  
The manifold hypothesis

# EfficientNet

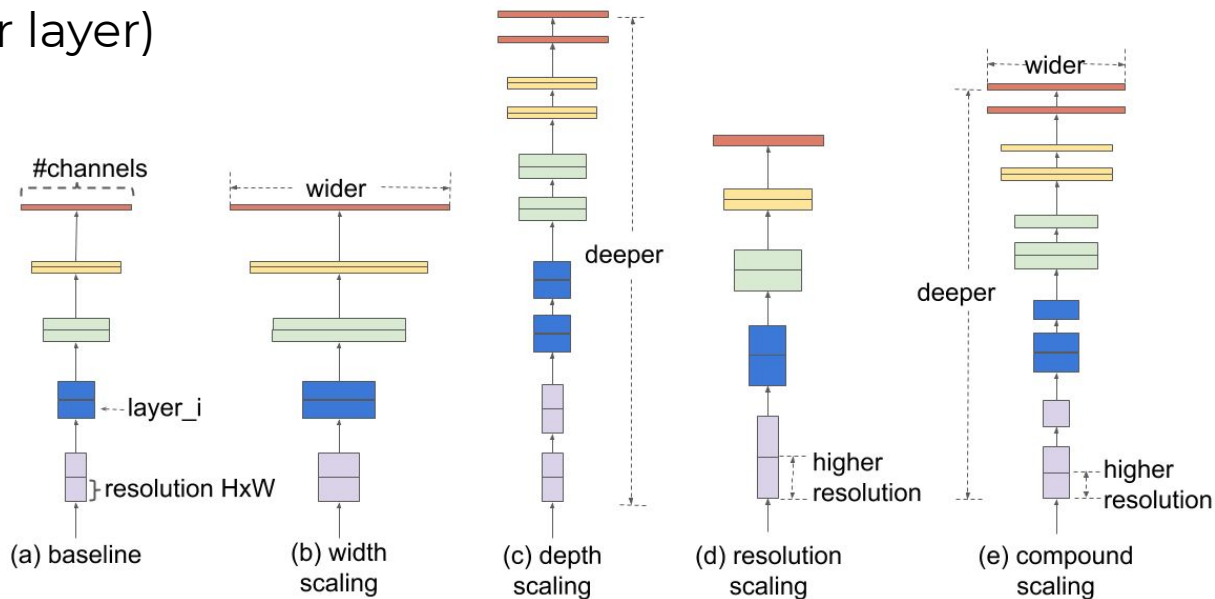
Should I go deeper, wider or bigger?

❖ Find a balance between them (they are all related!)

- Width (neurons per layer)
- Depth (layers)
- Resolution (input)

❖ Choose a size

- EfficientNetB0-B7



# Noisy Student

A semi-supervised training paradigm

1. Train model A (teacher) with the labeled data
  2. Use A to generate pseudo-labels for an unlabeled data set
  3. Train model B (student) with both labeled and pseudo-labeled data
- ❖ Iterate, re-labeling the unlabeled data each time
  - ❖ Highly regularized (noise!) student to guarantee improvement
  - ❖ Each student has more capacity than the previous





# Visualizing CNNs

# The Basics

- ❖ NN are representation learning techniques
- ❖ CNNs build hierarchically complex features
  - From Gabor filters to dog faces
  - Induced by convolution
  - Tend to focus on the “non obvious for humans”
    - Backgrounds, textures
- ❖ The closer to the loss, more classifier (task) and less representation (data)

# Ways of Looking at CNNs

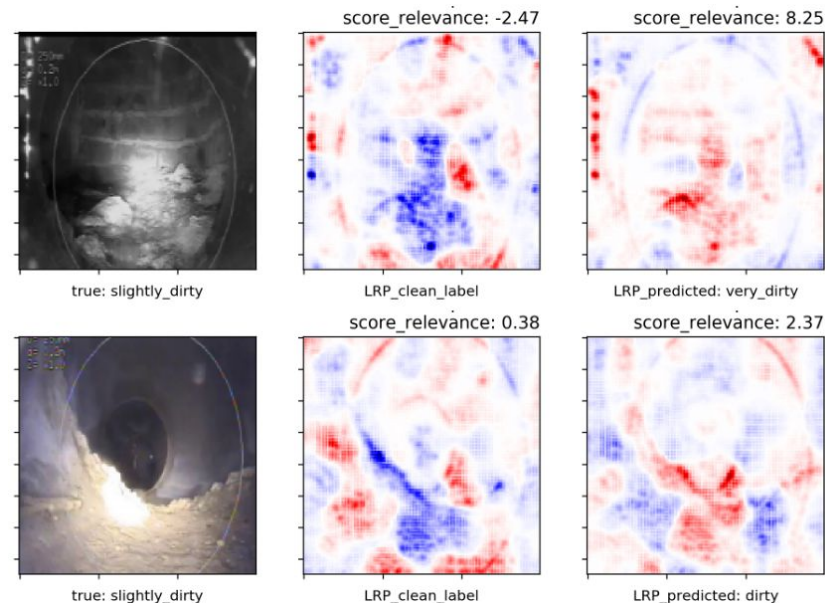

- ❖ *Attribution*: Where is the network looking?
  - Grounded. Instance based.
  - Explainability in practice.
- ❖ *Feature Visualization*: What is the network seeing?
  - Uncontextualized. Maximization based.
  - Diagnosys & Insight
- ❖ *Exemplification*: Which images cause a maximum activation?
  - Samples from a distribution





# Attribution

- ❖ Finding the importance of pixels
- ❖ Layerwise Relevance Propagation (LRP)
  - Backpropagate *an* output. Find the relevance of each neuron
    - Weighted by CNN parameters




# Feature Visualization

## ❖ Optimizing the input to maximize the output


- A neuron
- A channel
- A layer (DeepDream)



Low level




High level



# Exemplification

## ❖ Finding images within a dataset maximizing outputs

- Subjective
- Partial
- Stochastic






# Playing with CNNs

# Encoder-Decoder CNNs

- ❖ Pixel-wise classification task (image reconstruction loss)
- ❖ Bottlenecking makes it cheaper



● Conv + Batch Normalisation + ReLU

● Pooling

● Upsampling

● Softmax





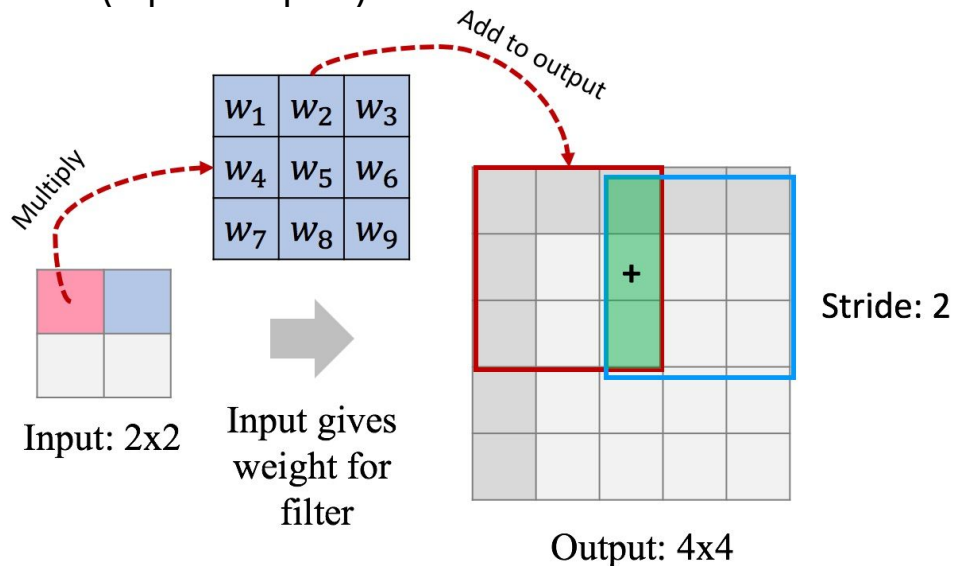
# Transposed Convolution ~~Deconvolution~~

❖ Reverse effect of regular convolution (upsample)

❖ Learnt interpolation

❖ Applications

- Segmentation
- GANs
- Super-Resolution
- Conv. Autoencoders



Input	Kernel					Output																																																										
<table border="1" style="display: inline-table; vertical-align: middle;"><tr><td>0</td><td>1</td></tr><tr><td>2</td><td>3</td></tr></table>	0	1	2	3	<table border="1" style="display: inline-table; vertical-align: middle;"><tr><td>0</td><td>1</td></tr><tr><td>2</td><td>3</td></tr></table>	0	1	2	3	=	<table border="1" style="display: inline-table; vertical-align: middle;"><tr><td>0</td><td>0</td><td></td></tr><tr><td>0</td><td>0</td><td></td></tr><tr><td></td><td></td><td></td></tr></table>	0	0		0	0					+	<table border="1" style="display: inline-table; vertical-align: middle;"><tr><td></td><td>0</td><td>1</td></tr><tr><td></td><td>2</td><td>3</td></tr><tr><td></td><td></td><td></td></tr></table>		0	1		2	3				+	<table border="1" style="display: inline-table; vertical-align: middle;"><tr><td></td><td></td><td></td></tr><tr><td>0</td><td>2</td><td></td></tr><tr><td>4</td><td>6</td><td></td></tr></table>				0	2		4	6		+	<table border="1" style="display: inline-table; vertical-align: middle;"><tr><td></td><td></td><td></td></tr><tr><td></td><td>0</td><td>3</td></tr><tr><td></td><td>6</td><td>9</td></tr></table>					0	3		6	9	=	<table border="1" style="display: inline-table; vertical-align: middle;"><tr><td>0</td><td>0</td><td>1</td></tr><tr><td>0</td><td>4</td><td>6</td></tr><tr><td>4</td><td>12</td><td>9</td></tr></table>	0	0	1	0	4	6	4	12	9
0	1																																																															
2	3																																																															
0	1																																																															
2	3																																																															
0	0																																																															
0	0																																																															
	0	1																																																														
	2	3																																																														
0	2																																																															
4	6																																																															
	0	3																																																														
	6	9																																																														
0	0	1																																																														
0	4	6																																																														
4	12	9																																																														

# Faster Segmentation

- ❖ ~~Pixel-wise classification~~ Object detection (bounding box)
  - Can be done with a “regular” CNN
- ❖ R-CNN: Propose crops (SVM). Extract features (CNN). Classify crops (SVM)
- ❖ Fast R-CNN: Extract features. Propose crops. Classify/Bounding Box (CNN)
- ❖ Faster R-CNN: Propose crops through a specific sub-net (RPN)
- ❖ YOLO v? (no regions, faster, less accurate)
  - Divide into grid. Predict class and bounding box for each cell.





# Better Segmentation

- ❖ Mask R-CNN
  - Faster R-CNN for object detection
  - FCN for instance segmentation (pixel classification)
- ❖ Xception
  - Depth-wise separable Convs (inverted order & w/o non-linearity)
  - Skip connections
  - Atrous SPP



# Style Transfer


- ❖ What do the correlation of activations intra-layer tell us?
  - What if we force it on another image?

- ❖ Gram matrix represents the *style*

- Channel-wise ( $c \times c$ )
- Several mid layers

- ❖ Activations represents the *content*

- One mid layer



- ❖ Optimize the **input** to minimize 2 losses
- ❖ Use a pre-trained net frozen
- ❖ Improved and extended

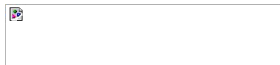
# Image Generation

- ❖ StyleGAN2 + pix2pixHD
  - Pixel-wise generative models
- ❖ Flow-edge Guided Video Completion
- ❖ <https://colab.research.google.com/drive/1KznIbRyNdiNBrrVbD7uolccdf9rngVUE?usp=sharing>



# Handwritten Generation

- ❖ <https://github.com/sjvasquez/handwriting-synthesis>
- ❖ <https://arxiv.org/abs/1308.0850>
- ❖ <https://www.calligrapher.ai/>



# References

- [1] [http://vordenker.de/ggphilosophy/mcculloch\\_a-logical-calculus.pdf](http://vordenker.de/ggphilosophy/mcculloch_a-logical-calculus.pdf)
- [2] <http://www-public.tem-tsp.eu/~gibson/Teaching/Teaching-ReadingMaterial/Rosenblatt58.pdf>
- [3] <http://www.dtic.mil/dtic/tr/fulltext/u2/236965.pdf>
- [4] [https://en.wikipedia.org/wiki/Perceptrons\\_\(book\)](https://en.wikipedia.org/wiki/Perceptrons_(book))
- [5] <http://www.andreykurenkov.com/writing/a-brief-history-of-neural-nets-and-deep-learning/>
- [6] [https://en.wikipedia.org/wiki/Perceptrons\\_\(book\)](https://en.wikipedia.org/wiki/Perceptrons_(book))
- [7] Werbos et al. "Beyond regression:" new tools for prediction and analysis in the behavioral sciences." Ph. D. dissertation, Harvard University (1974).
- [8] Rumelhart et al. "Learning Internal Representations by Error Propagation". MIT Press (1986).

# References

- [9] <https://towardsdatascience.com/effect-of-gradient-descent-optimizers-on-neural-net-training-d44678d27060>
- [10] <https://arxiv.org/abs/1711.05101>
- [11] <https://bbabenko.github.io/weight-decay/>
- [12] <https://towardsdatascience.com/weight-decay-l2-regularization-90a9e17713cd>
- [13] Veit, Andreas, Michael J. Wilber, and Serge Belongie. "Residual networks behave like ensembles of relatively shallow networks." Advances in neural information processing systems. 2016.
- [14] <https://thegradient.pub/semantic-segmentation/>
- [15] <https://arxiv.org/pdf/1603.08511>
- [16] <https://pdfs.semanticscholar.org/5c6a/0a8d993edf86846ac7c6be335fba244a59f8.pdf>



# References

- [17] <https://arxiv.org/pdf/1606.00915.pdf>
- [18] <https://arxiv.org/pdf/1610.02357.pdf>
- [19] [https://www.cv-foundation.org/openaccess/content\\_cvpr\\_2016/papers/Gatys\\_Image\\_Style\\_Transfer\\_CVPR\\_2016\\_paper.pdf](https://www.cv-foundation.org/openaccess/content_cvpr_2016/papers/Gatys_Image_Style_Transfer_CVPR_2016_paper.pdf)
- [20] <https://arxiv.org/abs/1603.08155>
- [21] <https://arxiv.org/abs/1603.03417>
- [22] <https://ai.googleblog.com/2016/10/supercharging-style-transfer.html>
- [23] <https://arxiv.org/pdf/1903.07291.pdf>
- [24] <http://nvidia-research-mingyuliu.com/gaugan>
- [25] Geirhos, Robert, et al. "ImageNet-trained CNNs are biased towards texture; increasing shape bias improves accuracy and robustness." arXiv preprint arXiv:1811.12231 (2018).

# References

- [26] Beery, Sara, Grant Van Horn, and Pietro Perona. "Recognition in terra incognita." Proceedings of the European Conference on Computer Vision (ECCV). 2018.
- [27] <https://distill.pub/2017/feature-visualization/>
- [28] <https://distill.pub/2018/building-blocks/>
- [29] Montavon, Grégoire, et al. "Layer-wise relevance propagation: an overview." Explainable AI: interpreting, explaining and visualizing deep learning. Springer, Cham, 2019. 193-209.
- [30] <https://medium.com/machine-intelligence-report/how-do-neural-networks-work-57d1ab5337ce>
- [31] Hebb, D.O. (1949), The organization of behavior, New York: Wiley



# References

- [32] Dauphin, Yann N., et al. "Identifying and attacking the saddle point problem in high-dimensional non-convex optimization." Advances in neural information processing systems. 2014.
- [33] Ruder, Sebastian. "An overview of gradient descent optimization algorithms." arXiv preprint arXiv:1609.04747 (2016).
- [34] Viazovetskyi, Yuri, Vladimir Ivashkin, and Evgeny Kashin. 'StyleGAN2 Distillation for Feed-Forward Image Manipulation'. 7 March 2020. <http://arxiv.org/abs/2003.03581>.
- [35] [https://medium.com/@jonathan\\_hui/gan-stylegan-stylegan2-479bdf256299](https://medium.com/@jonathan_hui/gan-stylegan-stylegan2-479bdf256299)
- [36] <https://www.justinpinkney.com/making-toonify/>
- [37] <http://chengao.vision/FGVC/files/FGVC.pdf>
- [38] Sandler, Mark, et al. "Mobilenetv2: Inverted residuals and linear bottlenecks." Proceedings of the IEEE conference on computer vision and pattern recognition. 2018.



# References

[39] Howard, Andrew G., et al. "Mobilenets: Efficient convolutional neural networks for mobile vision applications." arXiv preprint arXiv:1704.04861 (2017).



**Dario Garcia-Gasulla (BSC)**  
*[dario.garcia@bsc.es](mailto:dario.garcia@bsc.es)*

