

Deep Learning - MAI

Theory - RNNs

Dario Garcia Gasulla
dario.garcia@bsc.es

Context

Vanilla RNNs

Advanced RNNs

RNNs extensions



**Barcelona
Supercomputing
Center**

Centro Nacional de Supercomputación



Context

Dario Garcia Gasulla
dario.garcia@bsc.es

The Sequence

- ❖ Fully-connected and CNNs* have fixed inputs and outputs
- ❖ What if we have a variable shape input? Or output?
- ❖ Given a sequence (Text, video, audio, signal, bioinformatics...)
 - Learn the relations between the symbols that compose it

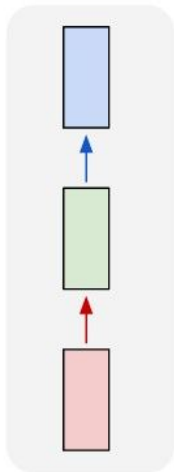
** CNNs can have a certain degree of variable inputs/outputs thanks to convolution*

What do we want

- ❖ Sequences can be processed using traditional methods (e.g., sliding windows), but sequences need to have the same length
- ❖ We want to...
 - process sequences of arbitrary length
 - provide different mappings (one to many, many to one, ...)

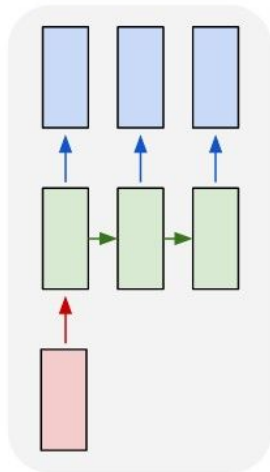
The ways of the sequence

one to one



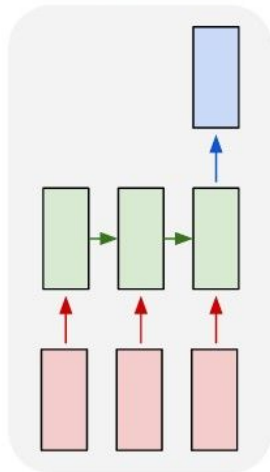
Not a RNN
(why?)

one to many



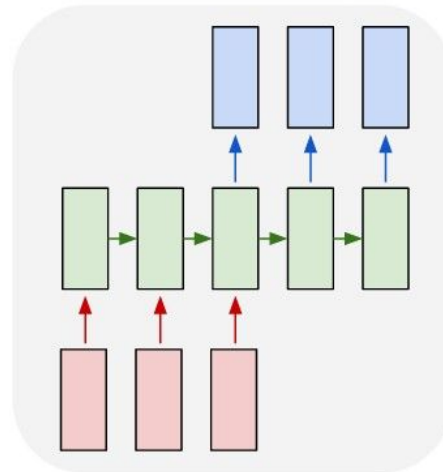
Sequence
output

many to one



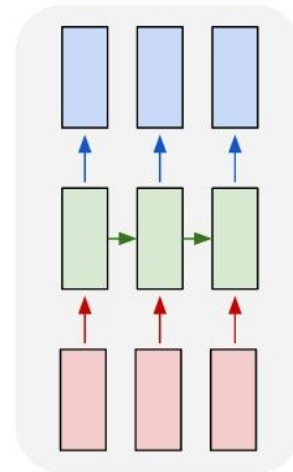
Sequence
input

many to many



Sequence to
sequence

many to many




Synced
sequence


One to many

❖ Image captioning

one to many



"man in black shirt is playing guitar."



"construction worker in orange safety vest is working on road."




"two young girls are playing with lego toy."

Many to one

❖ Sentiment analysis


- My flight was just delayed, s**t \Rightarrow Negative
- Never again BA, thanks for the dreadful flight \Rightarrow Negative
- We arrived on time, yeehaaa! \Rightarrow Positive
- Another day, another flight \Rightarrow Neutral
- Efficient, quick, delightful, always with BA \Rightarrow Positive

many to one



Sequence to Sequence


- ❖ Automatic translation
 - [How, many, programmers, for, changing, a,lightbulb,?] ⇒
 - [Wie, viele, Programmierer, zum, Wechseln, einer,Gl ühbirne,?] ⇒
 - [Combien, de, programmeurs, pour, changer, une,ampoule,?] ⇒
 - [¿,Cuantos, programadores, para, cambiar,una,bombilla,?] ⇒
 - [Zenbat, bonbilla, bat, aldatzeko,programatzaileak,?]



Synced Sequence

❖ Frame classification

many to many




Vanilla RNNs

Dario Garcia Gasulla
dario.garcia@bsc.es

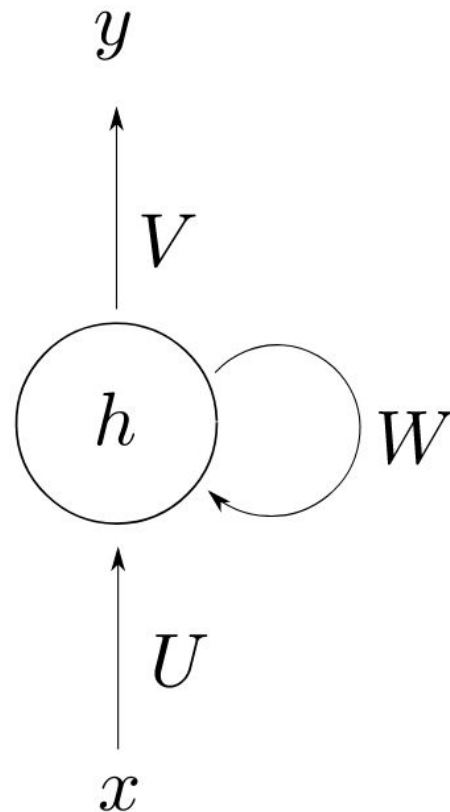
What makes a RNN?

- ❖ Recurrent Neural Networks are feed-forward networks with edges that span adjacent time steps (recurrent edges)
- ❖ On each step, a neuron receives inputs from data (as usual) and from previous time steps (history)
- ❖ In other words, previous time steps can influence future time steps
- ❖ RNNs are universal function approximators (Turing Complete)

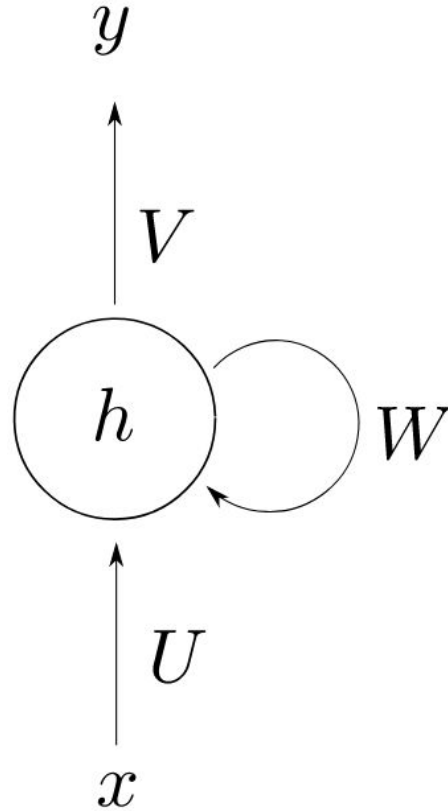


Who is who

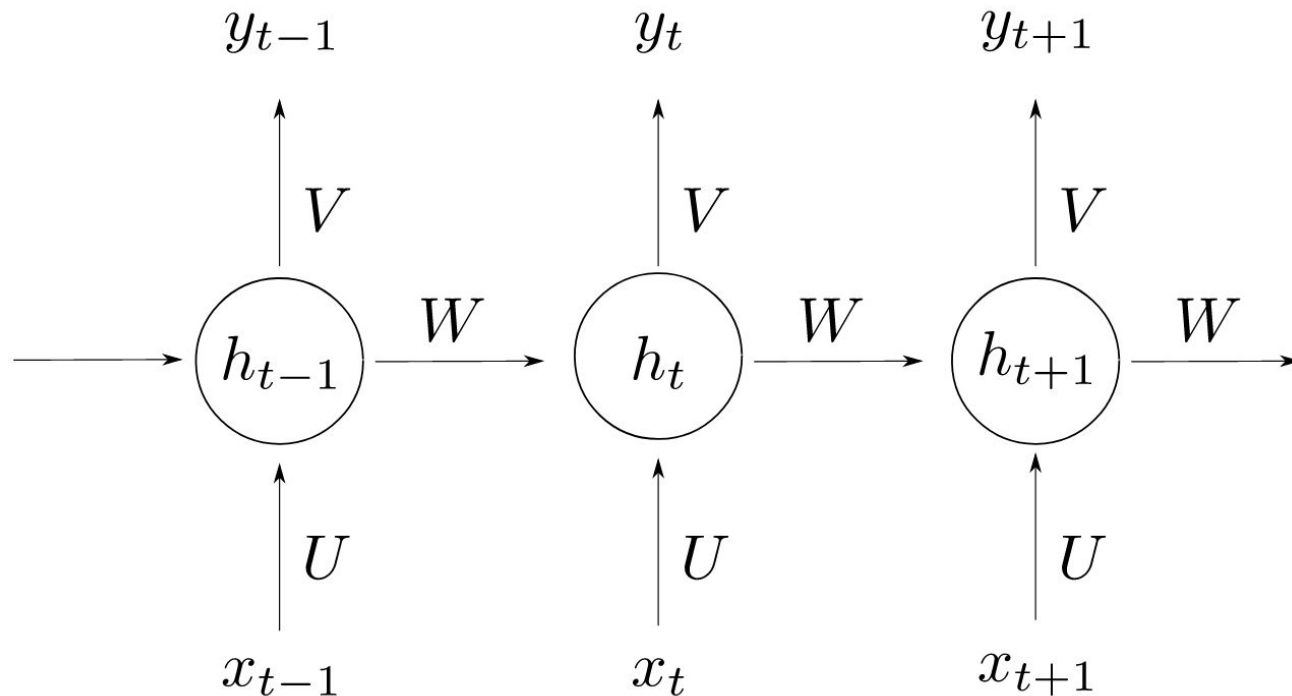
- ❖ RNNs Input (\mathbf{x}) is a vector of values for time t
- ❖ The hidden node (\mathbf{h}) stores the state
- ❖ Weights are shared through time (input independent!)
- ❖ Each step the computation uses the previous step
 - $h^{(t+1)} = f(h^{(t)}, x_{t+1}; \theta) = f(f(h^{(t-1)}, x_t; \theta), x_{t+1}; \theta) = \dots$
- ❖ RNN are a deep network that stacks layers through time



Before unrolling



After unrolling

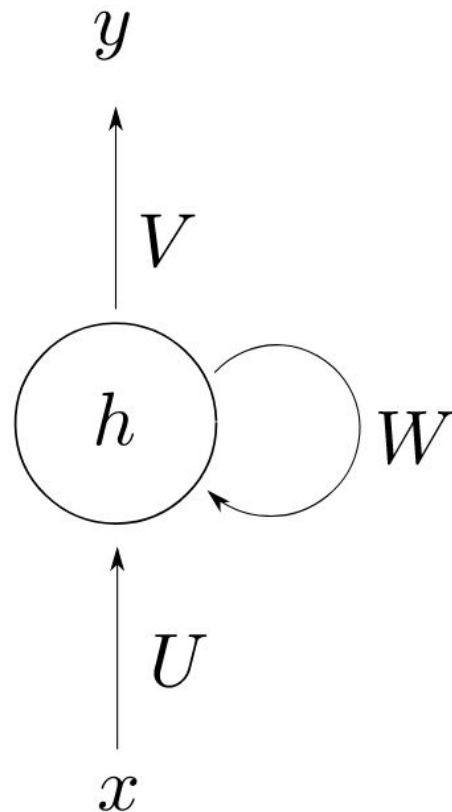


The computation

$$a^{(t)} = b + W \cdot h^{(t-1)} + U \cdot x^{(t)}$$

$$h^{(t)} = \tanh(a^{(t)})$$

$$y^{(t)} = c + V \cdot h^{(t)}$$




The power of RNNs lies in the unrolling, not in depth

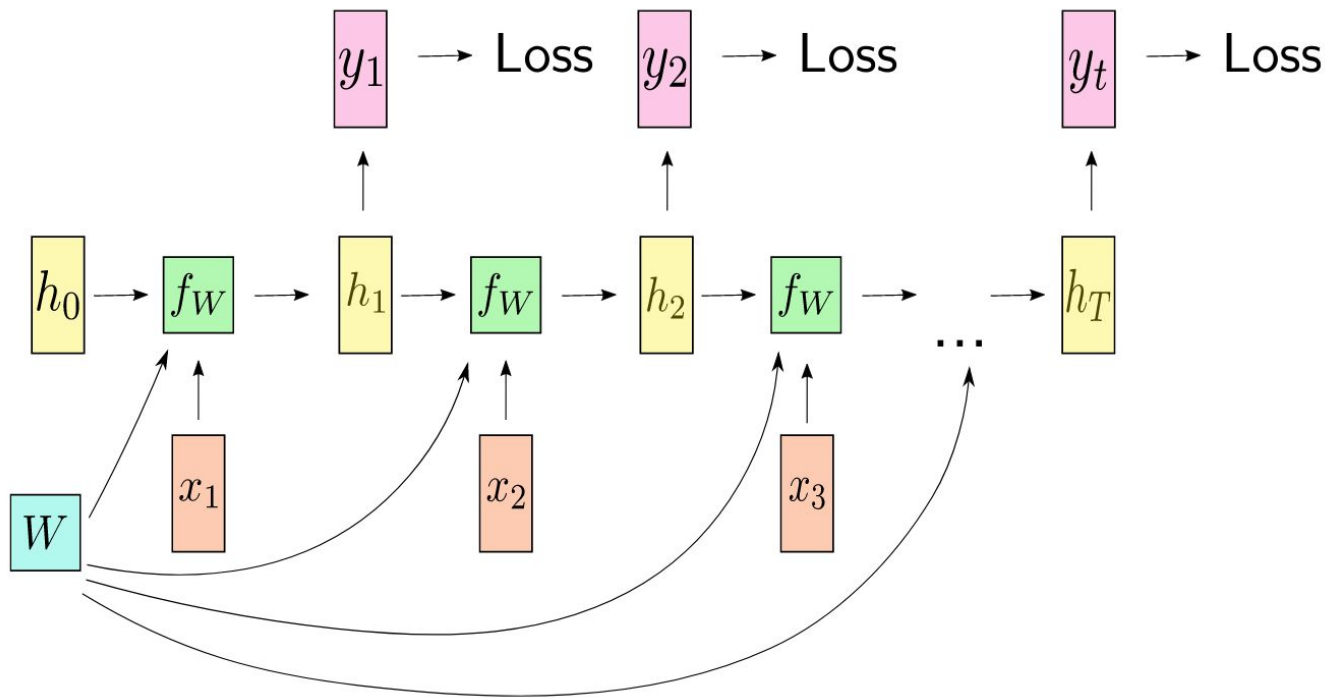
- RNNs with many layers are very expensive to compute
- Plus, the real challenge is in the memory
- Activation functions
 - Tanh is better than sigmoid
 - ReLU is also popular
- When working with text, inputs are most frequently *word embeddings*

How to train RNNs

- ❖ Backprop + SGD on the unfolded sequence
 - Compute activations and gradient
- ❖ Backpropagation Through Time (BPTT) by sum
- ❖ Input is limited in time to limit cost
 - Truncated BPTT
 - Influence limited to a time horizon



Combining losses



Training issues

- ❖ Sharing weights, with longer sequences, easily yields
 - Vanishing gradient (favours close by patterns wrt distant ones)
 - Exploding gradient
- ❖ Clipping gradients to prevent exploding
 - Scale gradient if the norm is above a threshold
- ❖ Vanishing gradients, next

Key features of RNNs

- ❖ Can process any length of input
 - For fixed sized inputs, consider other options
- ❖ Implements memory by design
- ❖ Model complexity is independent of input length
- ❖ All inputs processed with the same weights - reuse knowledge

RNNs inputs

- ❖ Word embeddings!
 - Language models
 - Unsupervised learning
 - Shallow NN
- ❖ word2vec vs glove (nobody cares anymore)
 - Expensive to train

Word embedding task

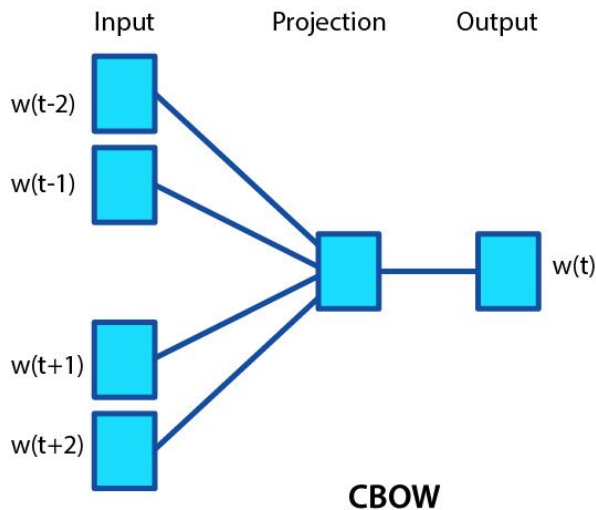
- ❖ Words are defined by their context
- ❖ Endless source of training data
- ❖ Use a sliding window of fixed length

Source Text	Training Samples						
<table><tr><td>The</td><td>quick</td><td>brown</td></tr></table> fox jumps over the lazy dog. ➡	The	quick	brown	(the, quick) (the, brown)			
The	quick	brown					
<table><tr><td>The</td><td>quick</td><td>brown</td><td>fox</td></tr></table> jumps over the lazy dog. ➡	The	quick	brown	fox	(quick, the) (quick, brown) (quick, fox)		
The	quick	brown	fox				
<table><tr><td>The</td><td>quick</td><td>brown</td><td>fox</td><td>jumps</td></tr></table> over the lazy dog. ➡	The	quick	brown	fox	jumps	(brown, the) (brown, quick) (brown, fox) (brown, jumps)	
The	quick	brown	fox	jumps			
<table><tr><td>The</td><td>quick</td><td>brown</td><td>fox</td><td>jumps</td><td>over</td></tr></table> the lazy dog. ➡	The	quick	brown	fox	jumps	over	(fox, quick) (fox, brown) (fox, jumps) (fox, over)
The	quick	brown	fox	jumps	over		

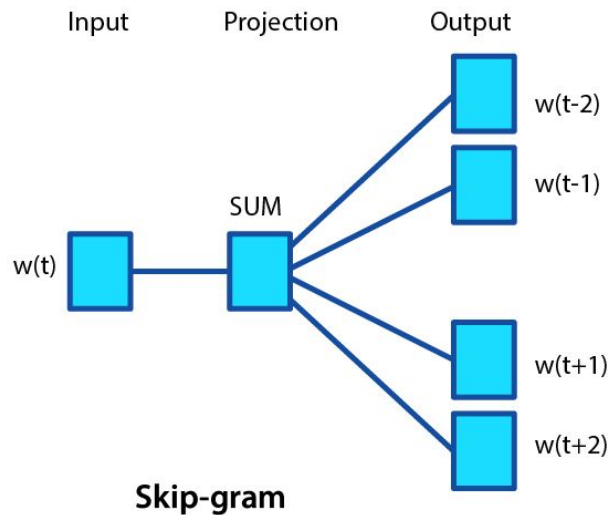
Word embedding models

Forget about the output!
Give me the intermediate representations

❖ Predict word given context



Predict context given word



❖ Faster and slightly better for frequent words

❖ Good for little training data and rare word representation

Word embedding properties

❖ Word similarity (cosine distance among vectors)

❖ Regularities

■ Add & Subtract

❖ So much bias...

❖ Go play

Relationship	Example 1	Example 2	Example 3
France - Paris	Italy: Rome	Japan: Tokyo	Florida: Tallahassee
big - bigger	small: larger	cold: colder	quick: quicker
Miami - Florida	Baltimore: Maryland	Dallas: Texas	Kona: Hawaii
Einstein - scientist	Messi: midfielder	Mozart: violinist	Picasso: painter
Sarkozy - France	Berlusconi: Italy	Merkel: Germany	Koizumi: Japan
copper - Cu	zinc: Zn	gold: Au	uranium: plutonium
Berlusconi - Silvio	Sarkozy: Nicolas	Putin: Medvedev	Obama: Barack
Microsoft - Windows	Google: Android	IBM: Linux	Apple: iPhone
Microsoft - Ballmer	Google: Yahoo	IBM: McNealy	Apple: Jobs
Japan - sushi	Germany: bratwurst	France: tapas	USA: pizza

■ <https://projector.tensorflow.org/>

<https://ronxin.github.io/wevi/>

Advanced RNNs

Dario Garcia Gasulla
dario.garcia@bsc.es

The limitations


- ❖ RNNs keep multiplying the same weight matrix over and over, which makes it error prone
- ❖ Keeping long term relations are hard for RNN
- ❖ The state has to encode both short and long term relations, which gets complicated as input sequences grow
- ❖ What if we make the state more powerful?

LSTM

- ❖ Long-Short Term Memory
 - In addition to the hidden state, add a cell state
 - Hidden state stores short-term info (large updates)
 - Cell state stores long-term info (small updates)
 - Include gate operators to erase, write and read from the cell
- ❖ Gates allow regulating the cell state, deciding the operation (e/r/w), its target (cell state segment) and magnitude (gate in range $[0,1]$), based on context

Inside a LSTM

- ❖ Forget gate: What is kept in the cell
 - $f_t = \sigma(W_f \cdot [h_{t-1}, x_t])$
- ❖ Input gate: What is added to the cell
 - $i_t = \sigma(W_i \cdot [h_{t-1}, x_t])$
 - $\hat{c}_t = \tanh(W_c \cdot [h_{t-1}, x_t])$ (new content)
- ❖ Output gate: What is passed to the hidden
 - $o_t = \sigma(W_o \cdot [h_{t-1}, x_t])$
 - $c_t = f_t c_{t-1} + i_t \times \hat{c}_t$ (new cell state)




$$h_t = o_t \times \tanh(c_t) \text{ (new hidden state)}$$

Why do LSTM work

- ❖ If forget gate is set to 0 (remember everything), long-term relations are always kept
- ❖ It includes a sort of short-cut, as long as you do not forget everything, gradient will flow!*

* No more W exp, but sigmoid can still vanish the gradient



Gated Recurrent Units

- ❖ A simplified version of LSTMs (thanks!)
 - No cell state
 - Update gate: What in the hidden state is updated/left as it is
 - LSTMs forget gate + input gate
 - Reset gate: Which part of the *previous* hidden state are used
 - Close to 1: Previous state has more relevance
 - Close to 0: New state has more relevance

Inside a GRU

- ❖ Update gate: How to change hidden state

- $z_t = \sigma(W_z \cdot [h_{t-1}, x_t])$

- ❖ Reset gate: How to combine

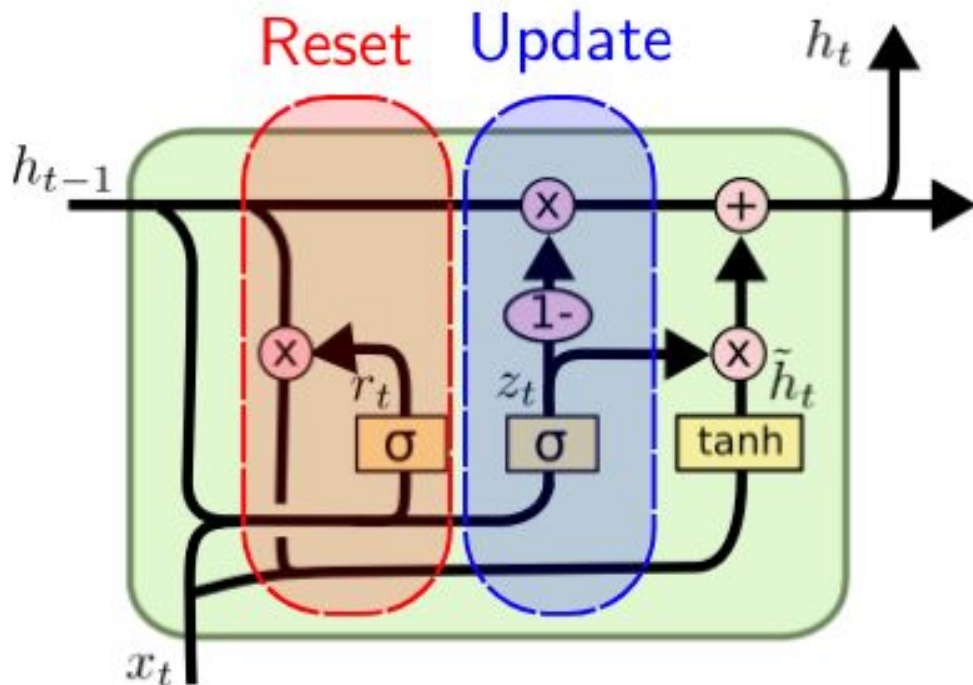
- $r_t = \sigma(W_r \cdot [h_{t-1}, x_t])$

- ❖ New hidden state content

- $\hat{h}_t = \tanh(W_h \cdot [r_t \times h_{t-1}, x_t])$

- ❖ Hidden state output

- $h_t = (1 - z_t) \times h_{t-1} + z_t \times \hat{h}_t$



Vanilla RNN vs LSTM vs GRU

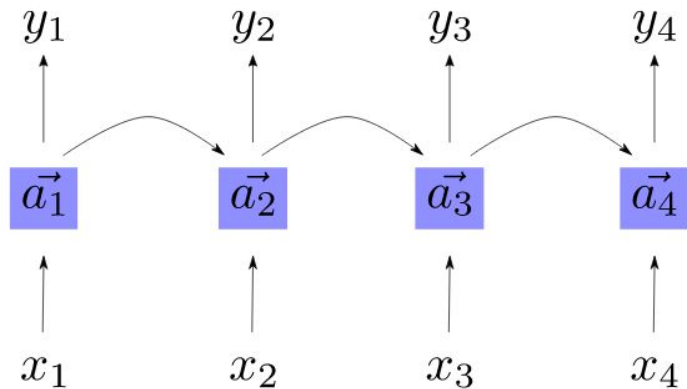
- ❖ Of all RNN variants, LSTMs and GRUs are the most widely used
 - Vanilla falls short in complex tasks, lacking long memory capacity
- ❖ Main difference between LSTM and GRUs: Complexity
 - Training GRUs is faster and includes less parameters
- ❖ Performance-wise, there are no consistent results

RNNs Extensions

Dario Garcia Gasulla
dario.garcia@bsc.es

Bidirectional RNNs


- ❖ Hidden states encode *past* states to influence current state



- ❖ What about future states? What if read from end to start?
 - PoS tagging, machine translation, speech/handwriting recognition

Bidirectional RNNs

- ❖ “Reading” in both directions at the same time
- ❖ Two RNNs, one in each direction, with different weights, concatenating their outputs




Training Bidirectional RNNs

- ❖ Both RNNs trained concurrently, but dependencies must be respected
 - Forward: Compute output of both RNNs and combine step by step
 - Backward: Compute gradient of both RNNs and combine step by step
- ❖ If possible (complete input sequence available), use bidirectionality

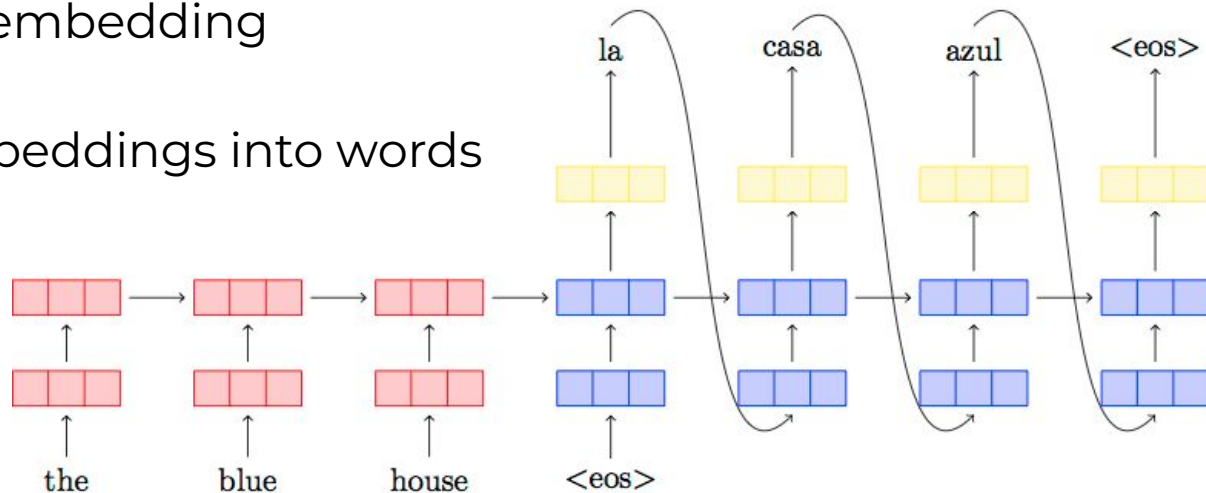
Multi-layer RNNs

- ❖ So far, only one layer used
 - Depth is provided by unrolling.
- ❖ Multi-layer RNNs
 - Stacking layers (rarely more than 4)
 - Provide extra abstraction capacity
 - A computational nightmare



Encoder-Decoder RNNs (seq2seq)

- ❖ Sequence to sequence of different length (Neural Machine Translation)
- ❖ One RNN encodes words within their context
- ❖ Generates a sentence embedding
- ❖ One RNN decodes embeddings into words
- ❖ Start and end tokens



Encoder-Decoder practical details

❖ Applications

- Automatic language translation
- Dialogue generation
- Document summarization
- Automatic response generation
- Input parsing

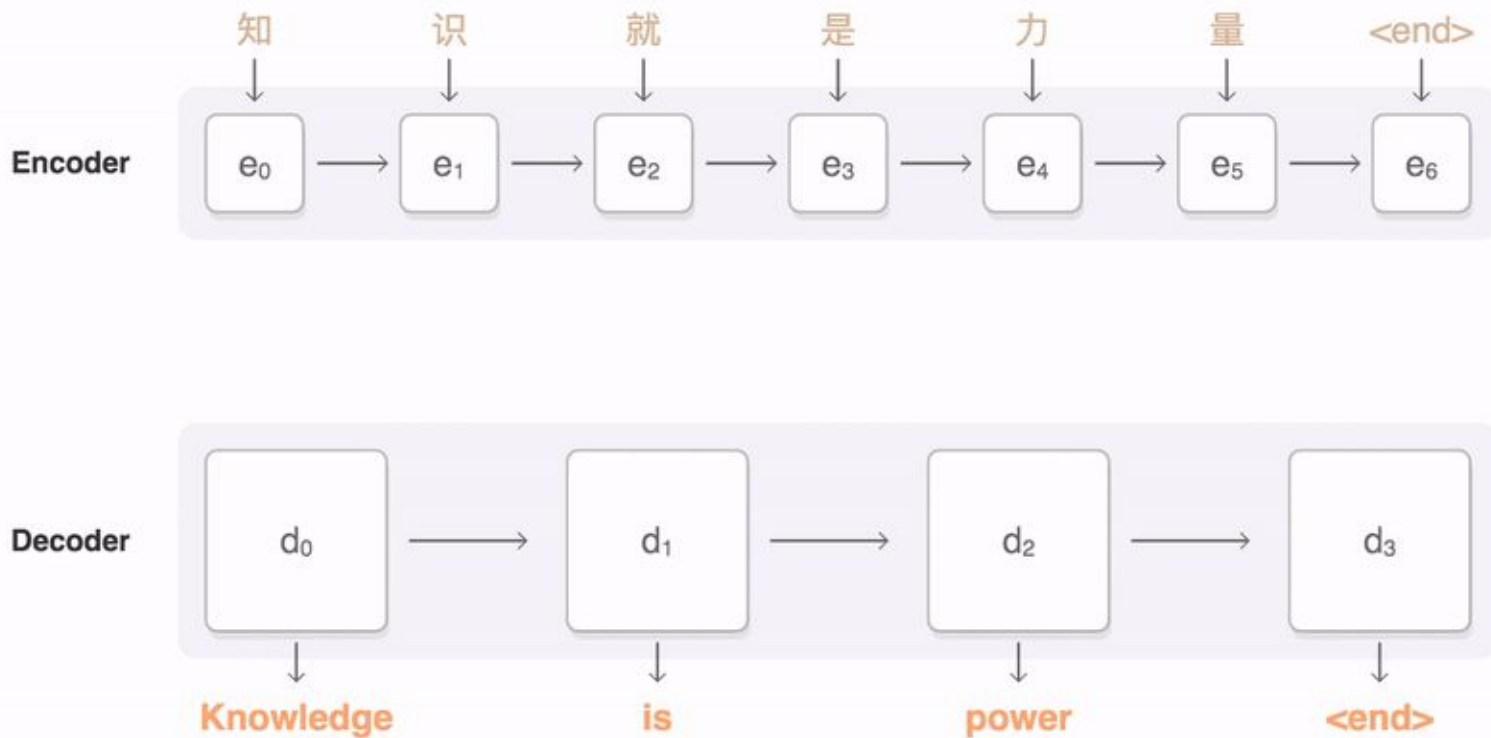
❖ Training

- Each decoding step generates one loss
- Negative log-likelihood of the true outcome at that point
- Average all losses at each step
- Decoder feedback replaced by true outcome

From Encoder-Decoder to Attention

- ❖ seq2seq limitations
 - Full sentence into a fixed-sized, unique embedding (bottleneck)
 - Different parts of the decoder focus on different parts of the input
 - Greedy decoding: Carrying on errors
 - Beam search decod.: Keep top k branches, find most probable path
- ❖ Solution: Attention
 - Let each decoder step decide which part of the input use


Attention overview



Seq2seq with attention


❖ Each decoder state

- Scores prev. hidden states (dot product)
- Turn into probabilities (softmax)
- Sum to make the **context vec.**
- Concatenate with hidden decoder state
- Output and fed to next step



Why seq2seq with attention

- ❖ Enables one different context for each decoding step
 - No fix-sized bottleneck
- ❖ Provides shortcuts (better gradient flows)
- ❖ More fine-grained -> better interpretability



A typical RNN model today...

“An encoder-decoder
bidirectional, multilayered
LSTM-based RNN with an
attention mechanism”

References

- [40] Andrej Karpathy, The Unreasonable Effectiveness of Recurrent Neural Networks, May 21, 2015. <http://karpathy.github.io/2015/05/21/rnn-effectiveness/>
- [41] <https://cs.stanford.edu/people/karpathy/deepimagesent/>
- [42] Hochreiter, Sepp, and Jürgen Schmidhuber. "Long short-term memory." Neural computation 9, no. 8 (1997): 1735-1780.
- [43] <http://colah.github.io/posts/2015-08-Understanding-LSTMs/>
- [44] Chung, Junyoung, Caglar Gulcehre, KyungHyun Cho, and Yoshua Bengio. "Empirical evaluation of gated recurrent neural networks on sequence modeling." arXiv preprint arXiv:1412.3555 (2014).
- [45] Jozefowicz, R., Zaremba, W., Sutskever, I. (2015). An empirical exploration of recurrent network architectures. In Proceedings of the 32nd International Conference on Machine Learning(ICML-15) (pp. 2342-2350).

References

[46] Chung, J., Gulcehre, C., Cho, K., Bengio, Y. (2014). Empirical evaluation of gated recurrent neural networks on sequence modeling. arXiv preprint arXiv:1412.3555.

[47]

https://calvinfeng.gitbook.io/machine-learning-notebook/supervised-learning/recurrent-neural-network/recurrent_neural_networks/

[48] Sutskever, Ilya, Oriol Vinyals, and Quoc V. Le. "Sequence to sequence learning with neural networks." arXiv preprint arXiv:1409.3215 (2014).

[49] Bahdanau, Dzmitry, Kyunghyun Cho, and Yoshua Bengio. "Neural machine translation by jointly learning to align and translate." arXiv preprint arXiv:1409.0473 (2014).

[50] Britz, Denny, Anna Goldie, Minh-Thang Luong, and Quoc Le. "Massive exploration of neural machine translation architectures." arXiv preprint arXiv:1703.03906 (2017).

References

- [51] Mikolov, Tomas, et al. Distributed Representations of Words and Phrases and their Compositionality.
- [52] Jeffrey Pennington, Richard Socher, and Christopher D. Manning. 2014. GloVe: Global Vectors for Word Representation.

Dario Garcia-Gasulla (BSC)
dario.garcia@bsc.es

