



# Deep Learning - MAI

Transfer Learning

**THEORY**

Dario Garcia Gasulla  
*dario.garcia@bsc.es*



# ***“Don’t be a hero”*** – Andrej Karpathy

The Transfer Learning philosophy

# Learning from scratch

- Trying to learn from scratch is difficult
  - You have to learn many things before getting to learn complex aspects of your task
- It's easier to learn if you have learnt something beforehand
  - There are some basic things needed to learn anything
  - Learning to “see”

# Why Transfer Learning


- You can learn faster
  - If I know that much, I'm that much closer to my goal
- You can learn better
  - There is a limited amount of things you can learn from data before getting trapped in spurious patterns.
  - What would you rather learn from your data?



# Putting things in perspective

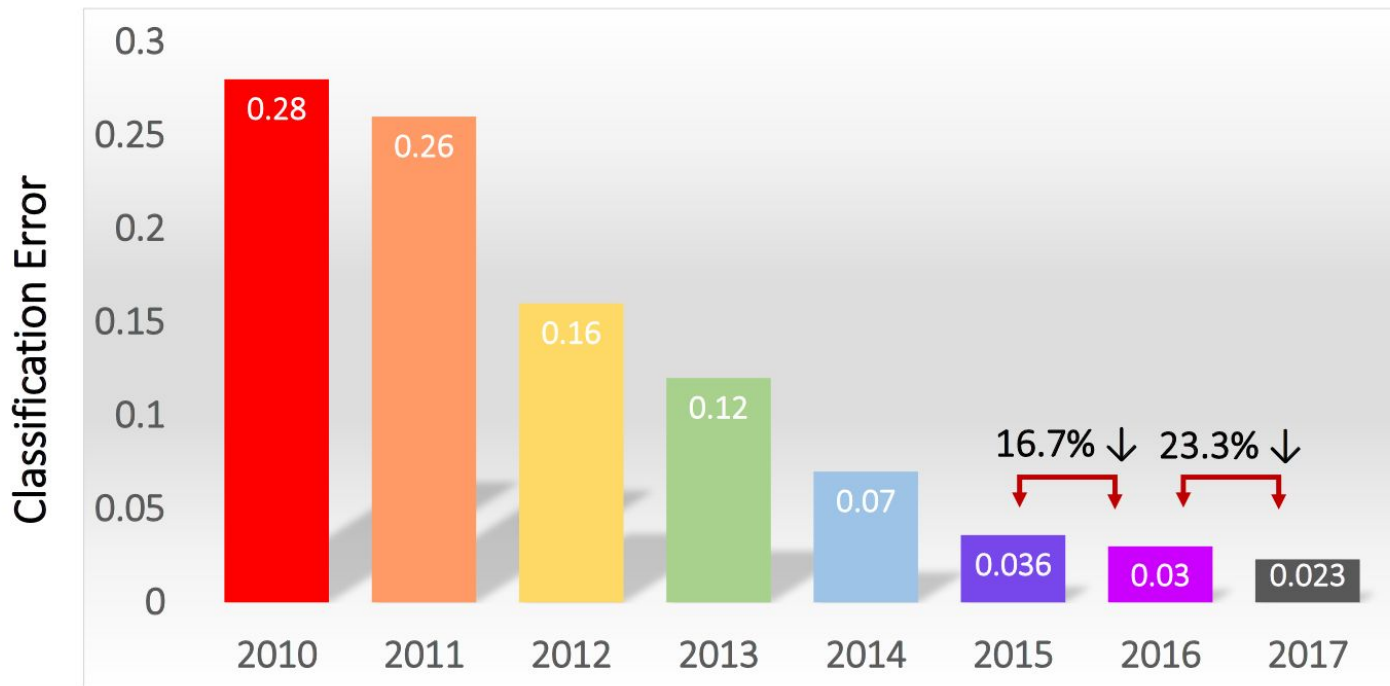
The ImageNet ¿success?

- 



# What we get

- We solved ImageNet



# What we pay

- Data labeling, transfer & storage
  - 1,000 images per class
- Training cost
  - Money (hardware, energy, salaries)
  - Environmental cost (CO2 emissions)
  - Human effort
    - Highly skilled professionals
    - Architecture design
    - Hyper-parameter fine tuning



# The ImageNet way is no way








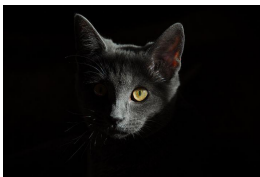

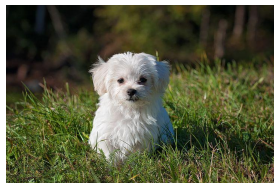
- We **cannot** do that for every single problem out there
  - The cost is too high. But more importantly...
- We **do not want to** do that for every single problem out there
  - TL to the rescue
- Transfer learning reduces the requirements on...
  - Data (implicit reuse of data)
  - Cost (reduced training costs)
  - Effort (halfway there)



# The essence of Transfer Learning

Learning it's all about generalization

# What is learning about?



# What is learning about?

Train set



Test set



# What is learning about









MODEL

PREDICT

Test set

Train set

TRAIN





# What is learning about

MODEL









PREDICT

Test set

Train set

TRAIN

FAIL



# What is the bias here?

Train set


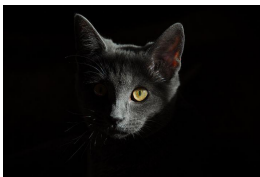



Test set




# What is the bias here?

Train set



GREEN  
BACKGROUND








DOG



# What is the bias here?

Train set



**GREEN  
BACKGROUND**



**DOG**

**Train and test sets  
have different  
conditional probability  
distributions**

# Fixing bias

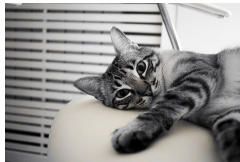

Train set



Solution?

# Fixing bias

Train set

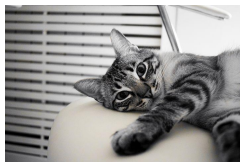
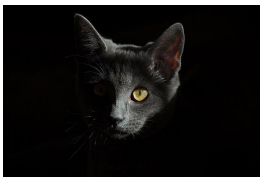
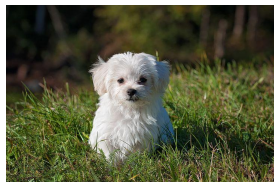


Solution?

Randomizing  
ensures that train and  
test sets have **similar**  
conditional probability  
distributions

# Fixing bias

Train set



Solution?


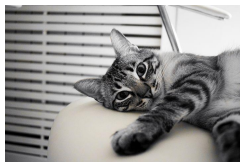
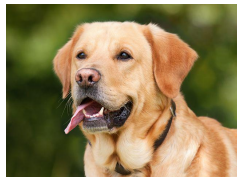
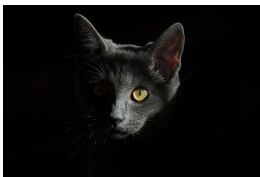
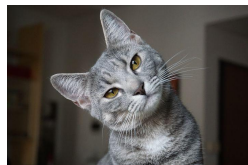
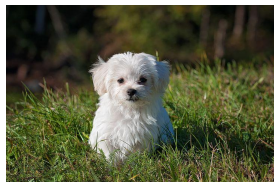
Randomizing

train and  
test  
data  
distribution  
they will never be  
exactly equal

comparing  
distributions

# Fixing bias

Train set



More similar means  
better generalization

Randomizing  
ensures that train and  
test sets have **similar**  
conditional probability  
distributions

# What is learning about?

Generalization between samples from the same source can be (approximately) ensured

**Train set**



**GENERALIZATION**






**Test set**





# What about?

Train set









Test set





# What about?

Generalization in this case is less certain  
Error is expected to rise  
Is it fixable?

Train set



Test set







# Formalizing Transfer Learning

## Tasks and Domains

Pan, Sinno Jialin, and Qiang Yang.  
"A survey on transfer learning."  
*IEEE Transactions on knowledge  
and data engineering* 22.10 (2010):  
1345-1359.

# Formalizing transfer learning


Domain:

Task:

# Formalizing transfer learning

**Domain:**  $\mathcal{D} = \{\mathcal{X}, P(X)\}$

- A feature space  $\mathcal{X}$




"The Elgar Concert  
Hall at the University  
of Birmingham for  
our third conference"

→ Bag of words

→ Content vector

- A marginal probability distribution  $P(X)$ , where  $X = \{x_1, \dots, x_n\} \in \mathcal{X}$




**Task:**

# Formalizing transfer learning

**Domain:**  $\mathcal{D} = \{\mathcal{X}, P(X)\}$

- A feature space  $\mathcal{X}$





$\neq$

"The Elgar Concert  
Hall at the University  
of Birmingham for  
our third conference"

→ Bag of words

→ Content vector

- A marginal probability distribution  $P(X)$ , where  $X = \{x_1, \dots, x_n\} \in \mathcal{X}$



$\neq$




**Task:**  $\mathcal{T} = \{y, f(\cdot)\}$

- A label space  $\mathcal{Y}$

**CAT, DOG  $\neq$  LION, WOLF**

- An objective predictive function  $f(\cdot) \Leftrightarrow P(y|x)$




# Formalizing transfer learning

**Domain:**  $\mathcal{D} = \{\mathcal{X}, P(X)\}$

- A feature space  $\mathcal{X}$ 
  - The Same (different)
- A marginal probability distribution  $P(X)$ 
  - Different
  - Similar

**Task:**  $\mathcal{T} = \{y, f(\cdot)\}$

Source



Target



# Formalizing transfer learning


**Domain:**  $\mathcal{D} = \{\mathcal{X}, P(X)\}$

- A feature space  $\mathcal{X}$ 
  - The Same (different)
- A marginal probability distribution  $P(X)$ 
  - Different
  - Similar

**Task:**  $\mathcal{T} = \{y, f(\cdot)\}$

- A label space  $y$ 
  - Different
  - The same
- An objective predictive function
  - Different (but similar?)

Source



{CAT, DOG}  
{FELINE, CANINE}

$f_S(\cdot)$

Target







{LION, WOLF}  
{FELINE, CANINE}

$f_T(\cdot)$

# What is **transfer** learning about?

**Train set**





**Test set**









# What is **transfer** learning about?

## Target domain






## Source domain







# What is **transfer** learning about?

Source domain



Source Task  $\psi = \{\text{CAT, DOG}\}, f_s(\cdot)$

Target domain



Target Task  $\psi = \{\text{LION, WOLF}\}, f_t(\cdot)$



# Formalizing transfer learning


**Domain:**  $\mathcal{D} = \{\mathcal{X}, P(X)\}$

- A feature space  $\mathcal{X}$ 
  - The Same (different)
- A marginal probability distribution  $P(X)$ 
  - Different
  - Similar

**Task:**  $\mathcal{T} = \{y, f(\cdot)\}$

- A label space  $y$ 
  - Different
  - The same

Source



{CAT, DOG}  
{FELINE, CANINE}

Target



{LION, WOLF}  
{FELINE, CANINE}

- An objective predictive function
  - Different (but similar?)

$f_S(\cdot)$

$f_T(\cdot)$



# Formalizing transfer learning

Source

$f_S(\cdot)$

Target

$f_T(\cdot)$



- Are they similar?
- Can we just use  $\hat{f}_S(\cdot)$  to approximate  $f_T(\cdot)$ ?
- Can we reuse part of it?

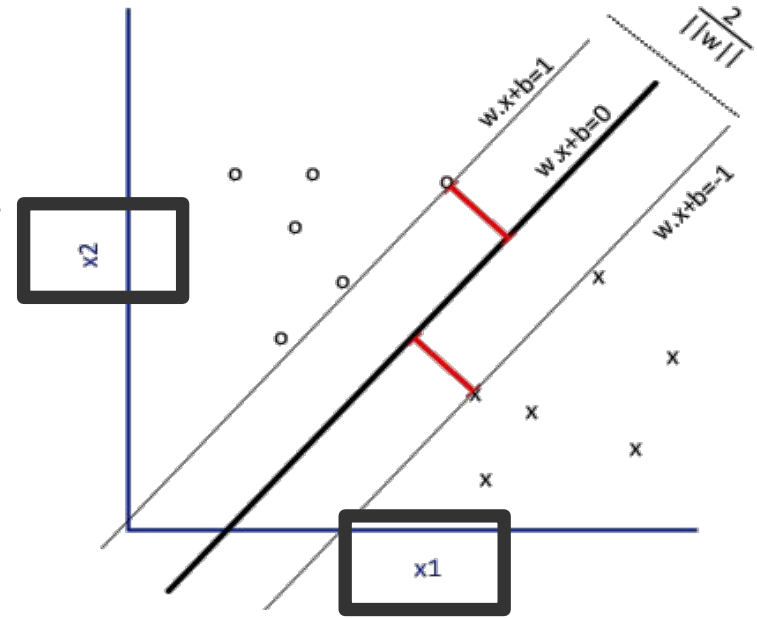


# Representation Learning & Classifiers

Learning to describe


# A typical classifier

- **Support Vector Machine (SVM)** is just a **classifier** (a very good one).
- SVM find the best boundary separating the data instances into different classes in a **given** feature space.

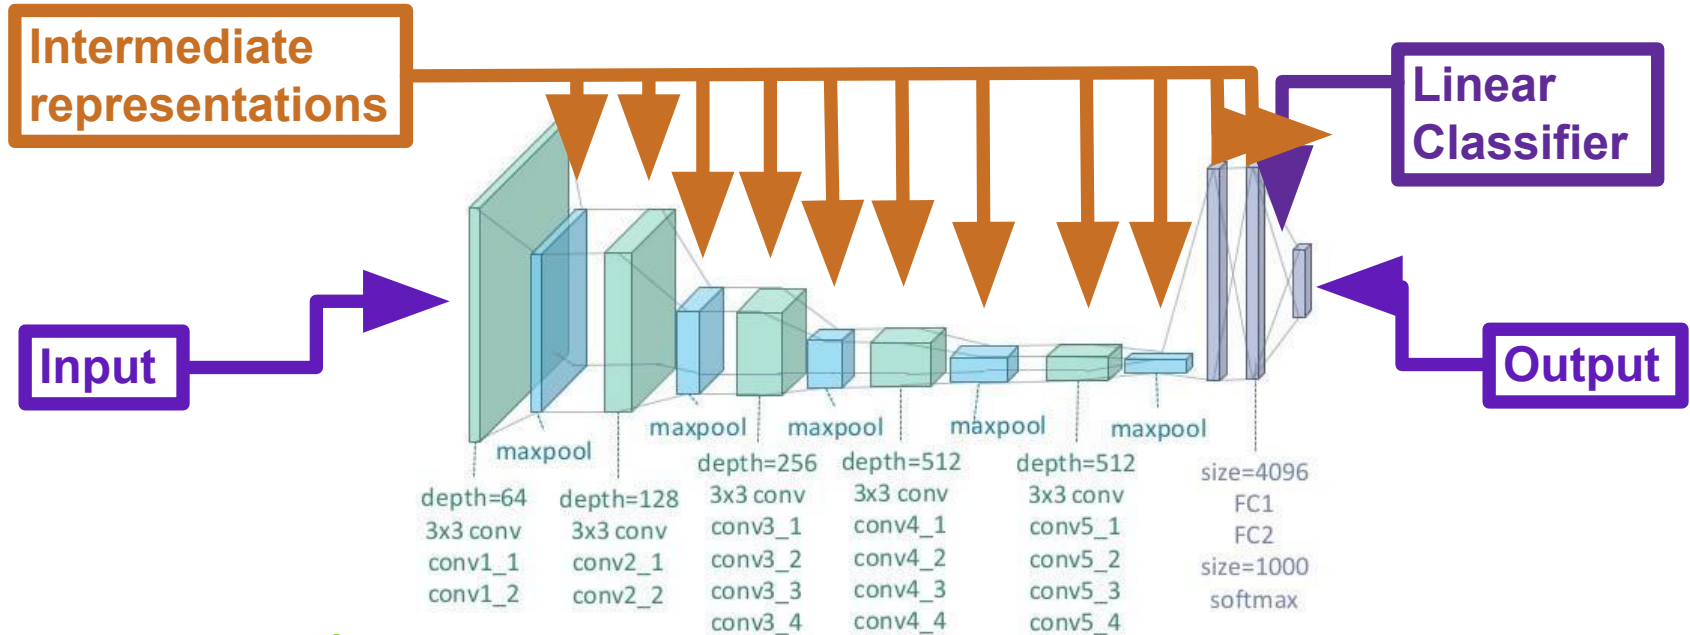


# A good classifier

- SVMs using the **kernel trick** can overcome the linear limitation through an **implicit** mapping to a higher dimensional feature space



# Deep Neural Networks and classifiers






# Reusing Deep Representations

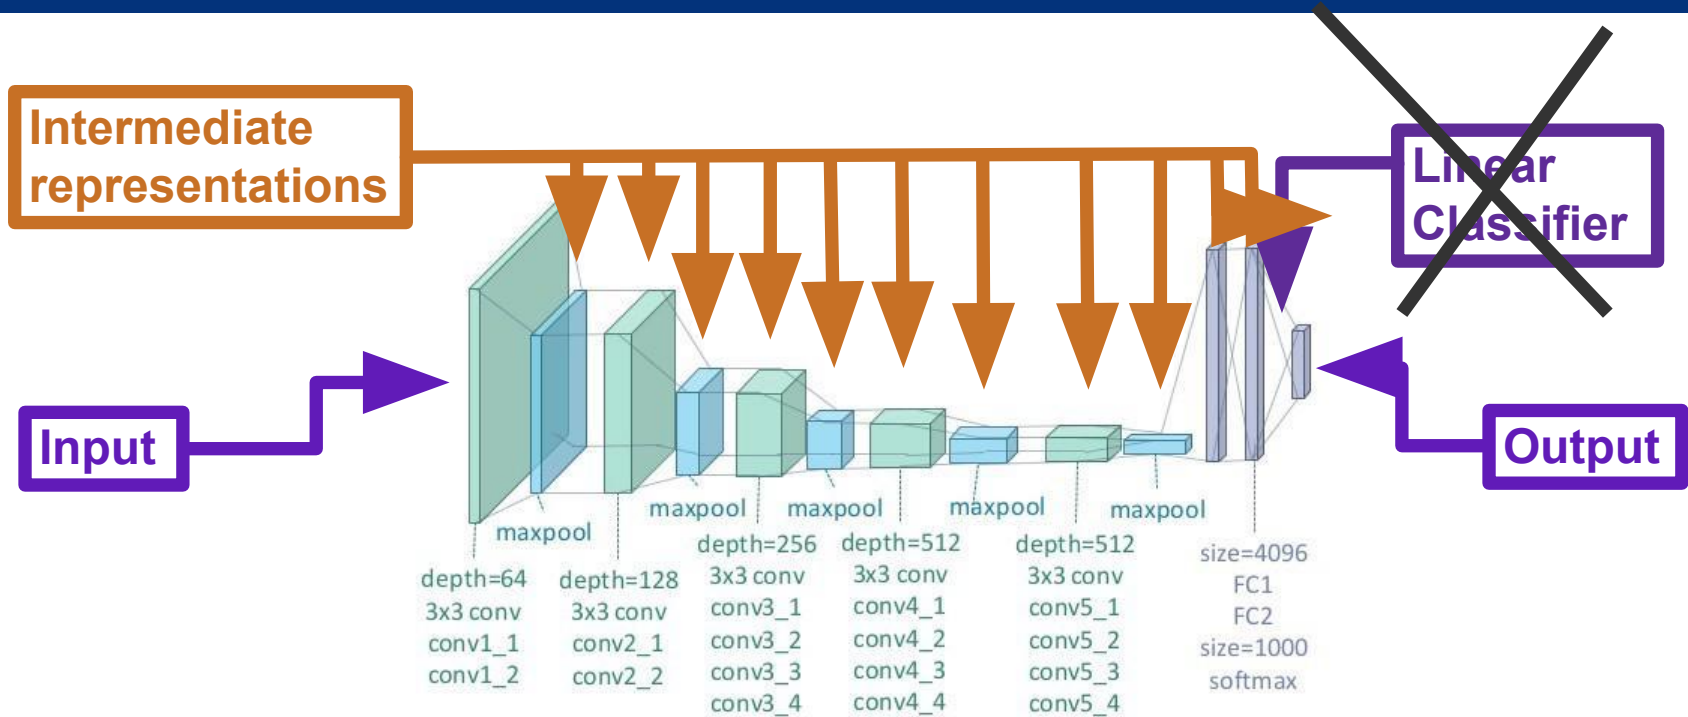
Save the Earth - Reuse DNNs



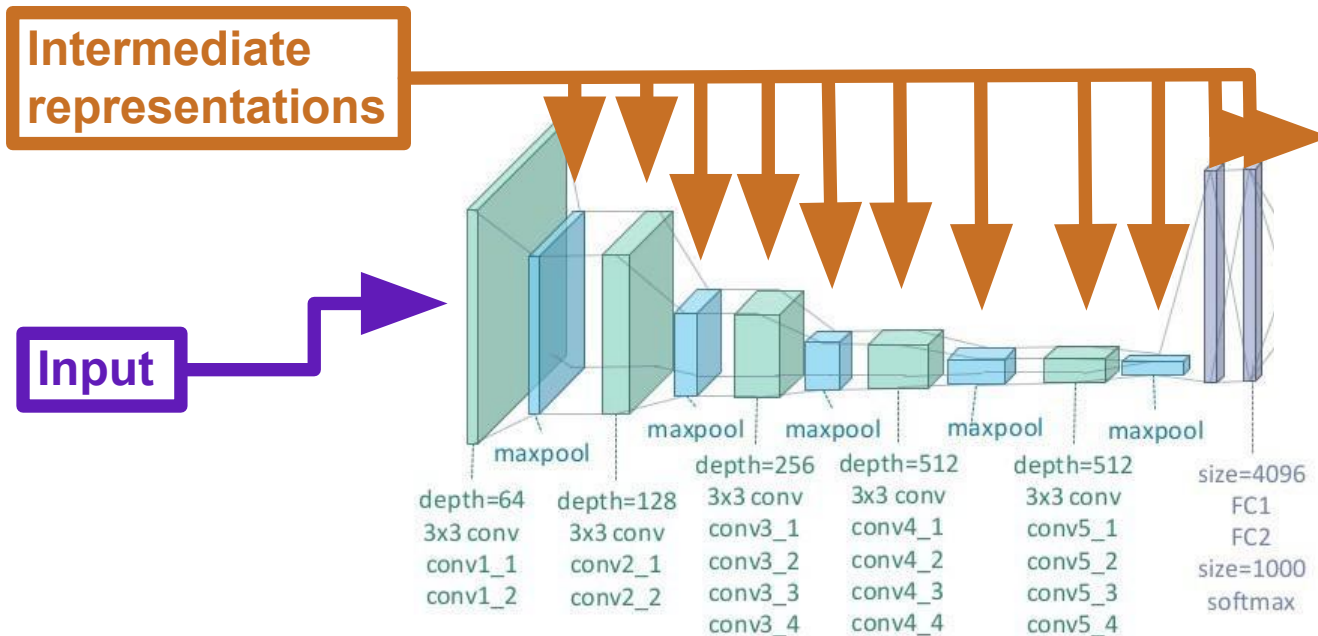
# What can be saved?




# What can be saved?



# Feature extraction




# Feature extraction



# Reuse All

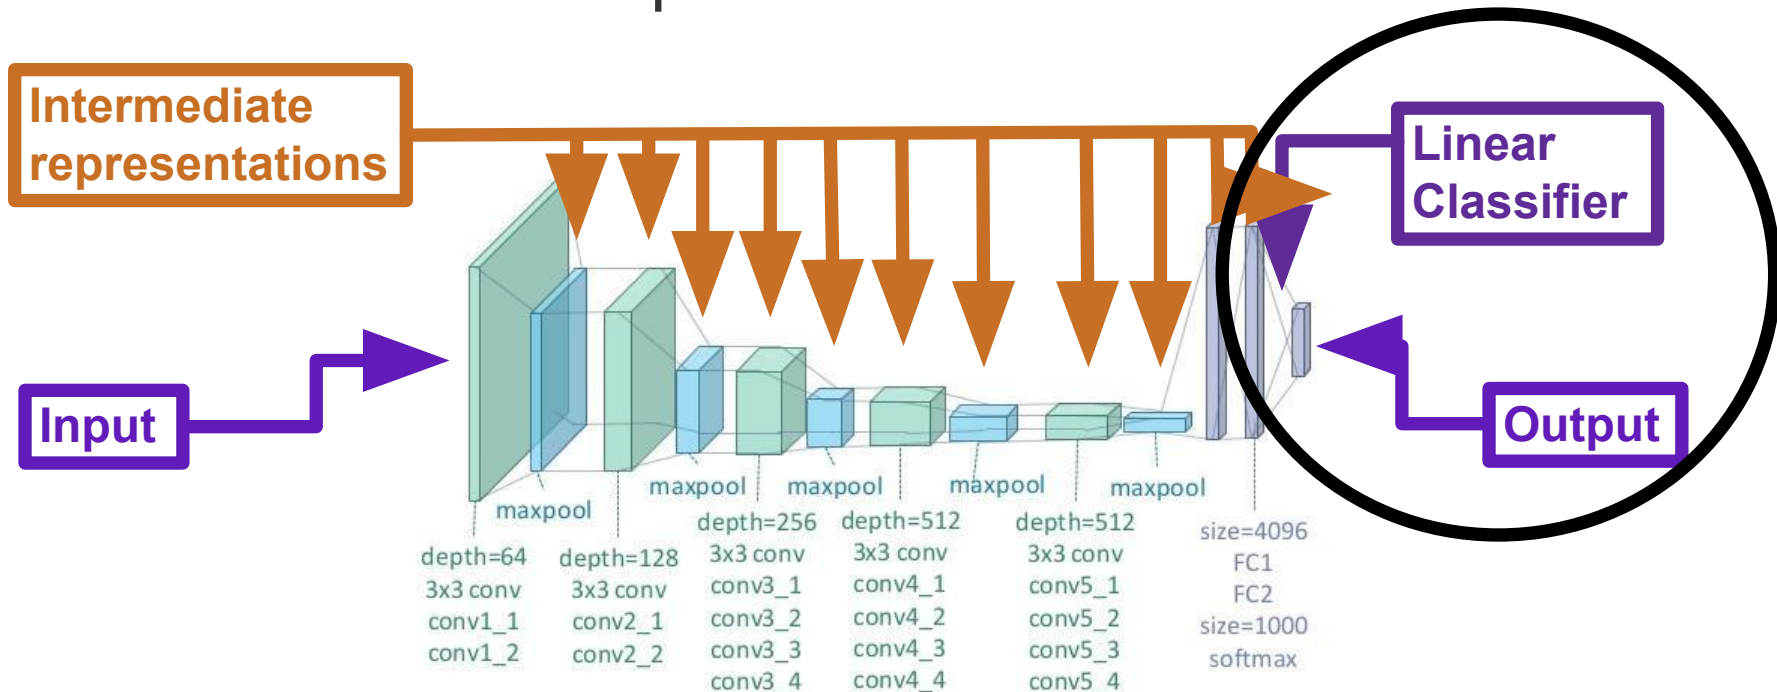
- ❖ DNN last layer features + SVM
  - Feature extraction
  - Similar task and domain



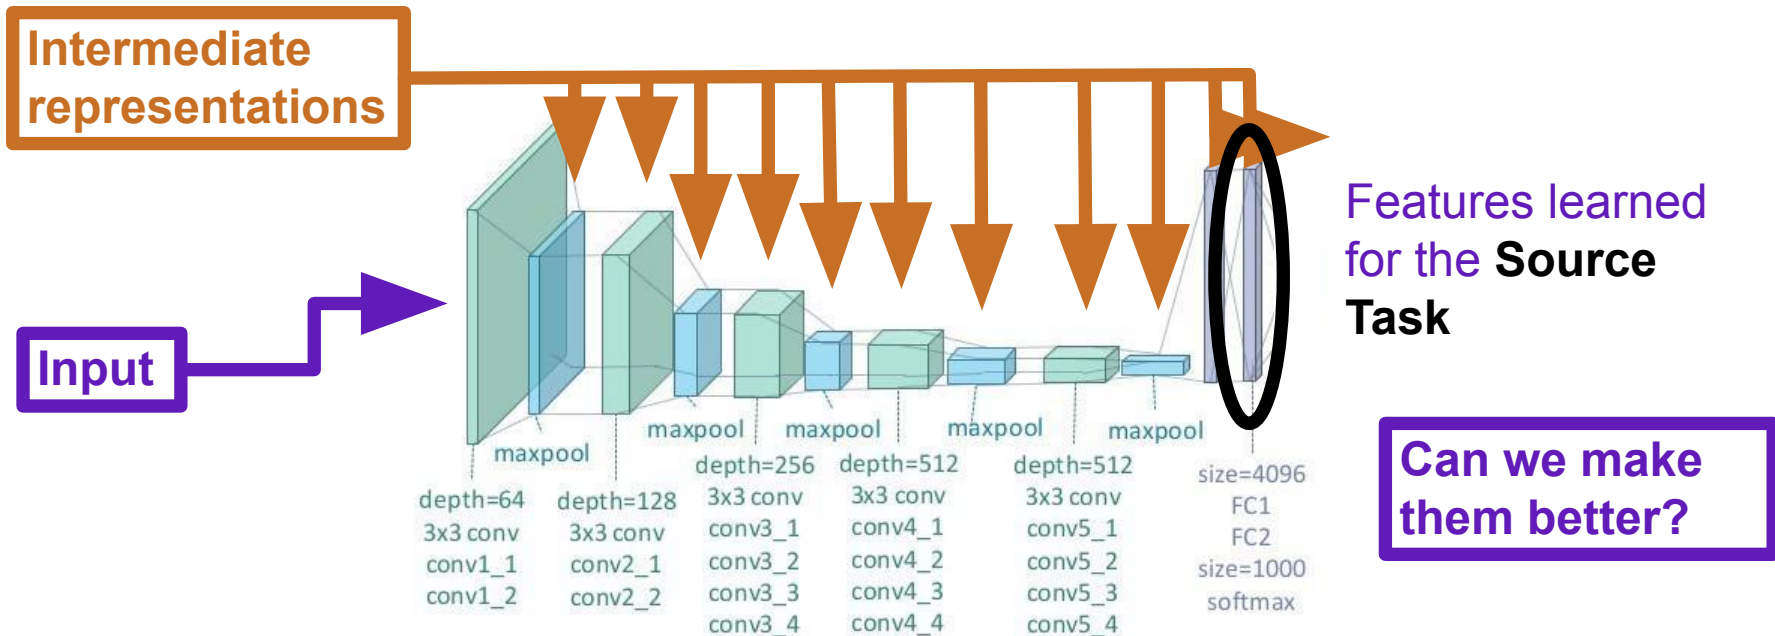
# Fine tuning

What if the tasks are quite different?

Source Task

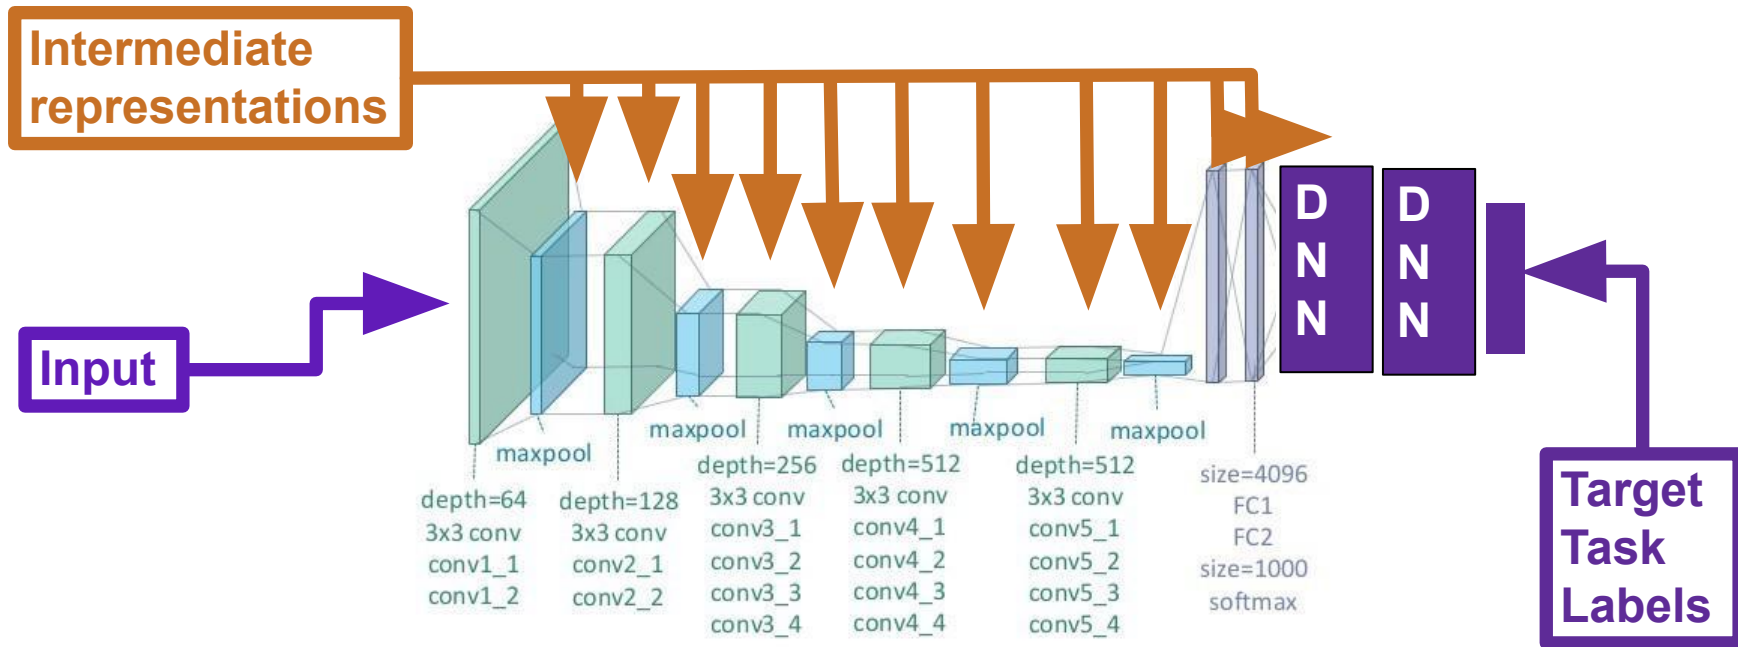


# Fine tuning






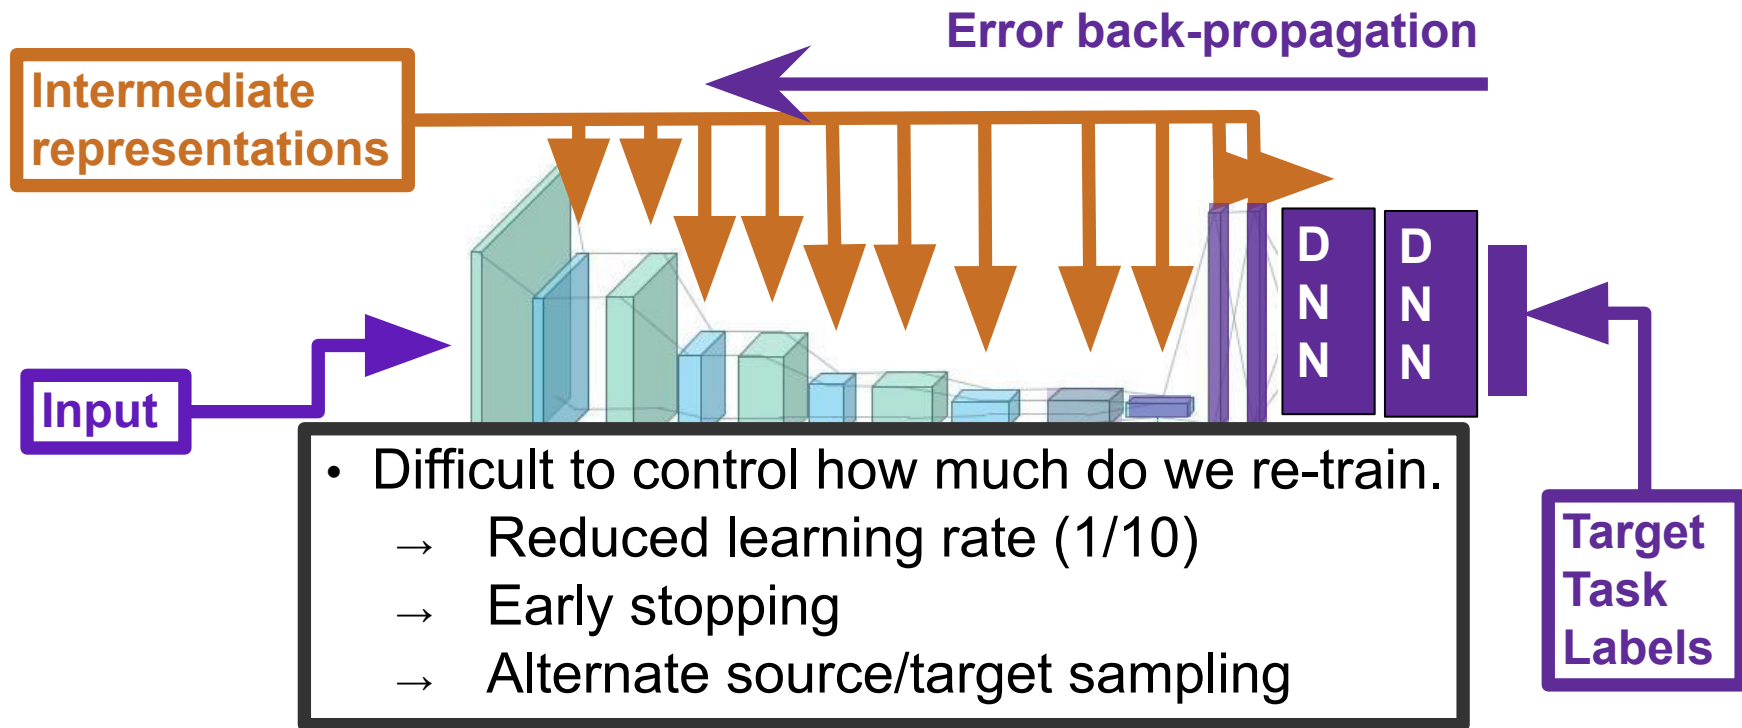
# Fine tuning



# Fine tuning





# Fine tuning

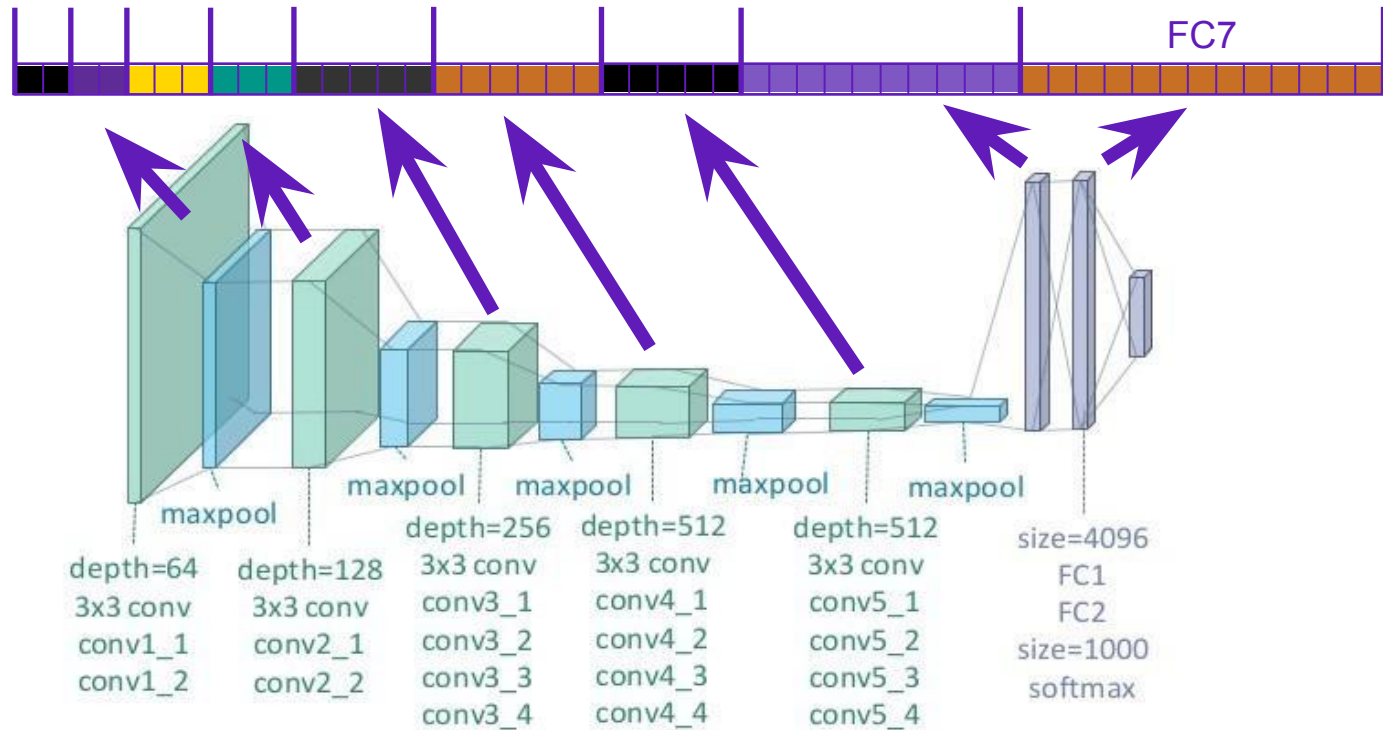


# Retrain


- ❖ DNN last layer features + SVM
  - Feature extraction
  - Similar task and domain
- ❖ Train one or several NN layers + pre-trained layers
  - Fine tuning
  - Data volume




# Knowledge inside DNN



# Knowledge inside DNN



# Knowledge inside DNN



Visualizations from: Yosinski, Jason, et al. "Understanding neural networks through deep visualization." *arXiv preprint arXiv:1506.06579* (2015).





# Fine Tuning

To improve, to remember, to forget

# The choices in fine tuning


- Reuse and **freeze**
  - Use source status
  - “Its good as it is”
- Reuse and **fine tune**
  - Start from source status, adjust with target
  - “Its a good starting point”
- **Random** init
  - Reinitialize weights randomly, train with target only
  - “Its pretty much useless”

# The order of fine tuning

**Freeze**

**Fine tune**


**Random**



# The order of fine tuning

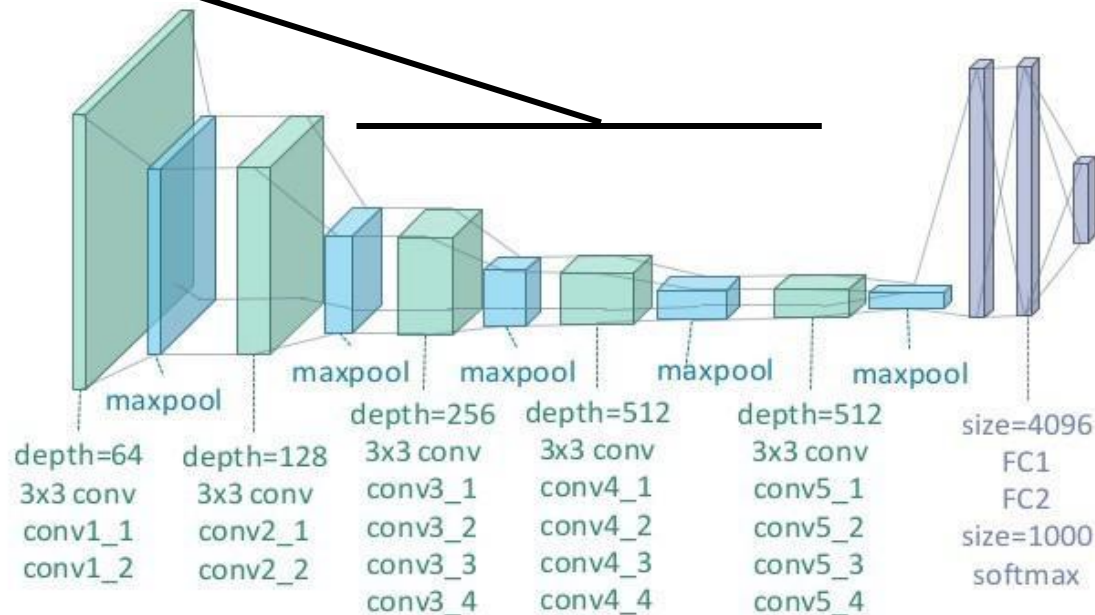
**Freeze**

Layers which are pretty stable and universal  
Increase with source-target similarity



# The order of fine tuning

**Fine tune** Layers which are pretty similar but improvable  
Increase with source-target similarity & data volume

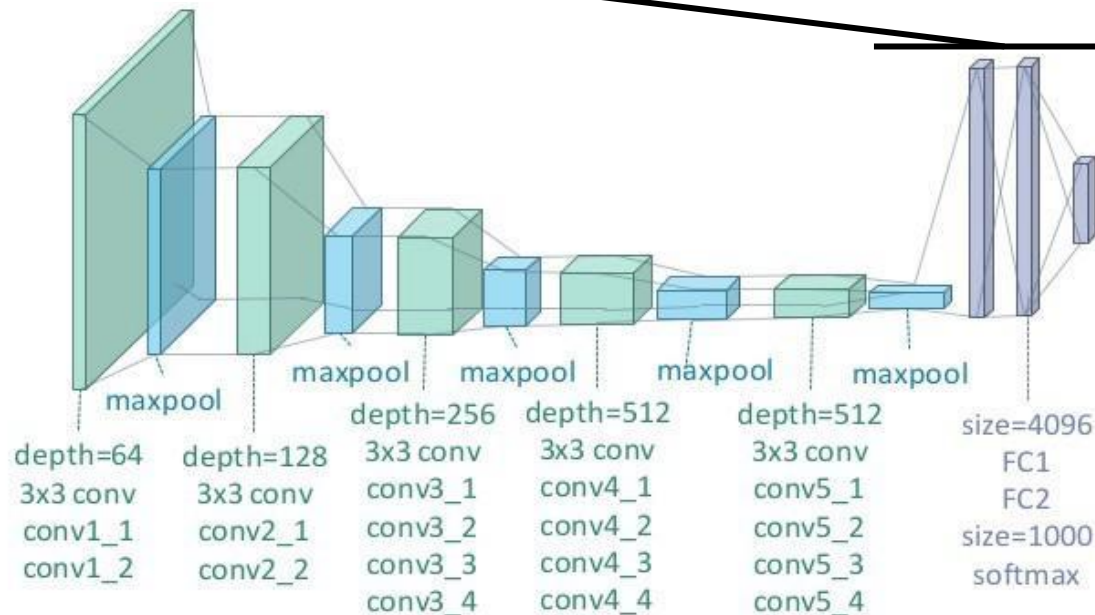


# The order of fine tuning

## Random

Layers which are pretty dissimilar

Increase with source-target dissimilarity & data volume



# Trade-off of fine tuning

- Reuse and **freeze**
  - Remove parameters for target to learn (needs data but allows focus)
  - Adds noise
- Reuse and **fine tune**
  - Allows to focus learning (requires data)
  - Adds bias
- **Random** init
  - Again, from the top (cost, cost, cost)
  - Tailor made for target





# Feature Extraction

To improve, to remember, to forget

# Factors deep representations quality

- Source task
  - Total volume
  - Class variety
- Target task
  - Source-target similarity
- Model
  - Capacity
  - Accuracy



# Factors deep representations quality


- Source task
  - Total volume
  - Class variety
- Target task
  - Source-target similarity
- Model
  - Capacity
  - Accuracy

If you have all of this,  
feature extraction plus a  
classifier will get you close  
to state-of-the-art




# Which layers to use?

- If source & target task are VERY similar use the “classifier” layers



# Which layers to use?


- If source & target task are NOT very similar  
broaden the scope



# Feature extraction normalization

- When doing feature extraction for a regular classifier (e.g., SVM) each feature is assumed to be i.i.d. (not even close!)
- Beware of size
  - FC layers have lots of activations
  - Conv layers activations are spatially dependent
- Beware of scale
  - Different layers activate with different strength
- Default solution: L2-norm (by layer)
  - Does not fix scale (careful if mixing layers!)

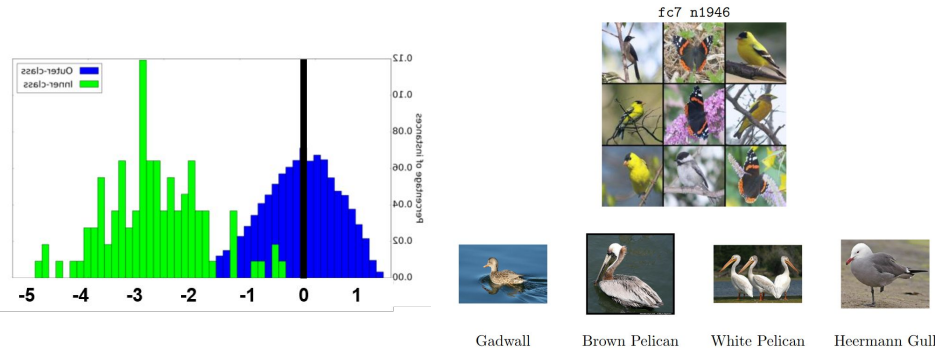

VGG16	Convs	FCs
# Layers:	14	2
Activations:	33%	66%



# Advanced feature normalization

- Normalizing features considering the target
  - Feature standardization (vertically instead of horizontally)
- For each feature..
  - Compute mean and std dev. on *target* training set
  - Normalize feature-wise to zero mean, one std dev.
  - Features are adapted to *target domain*

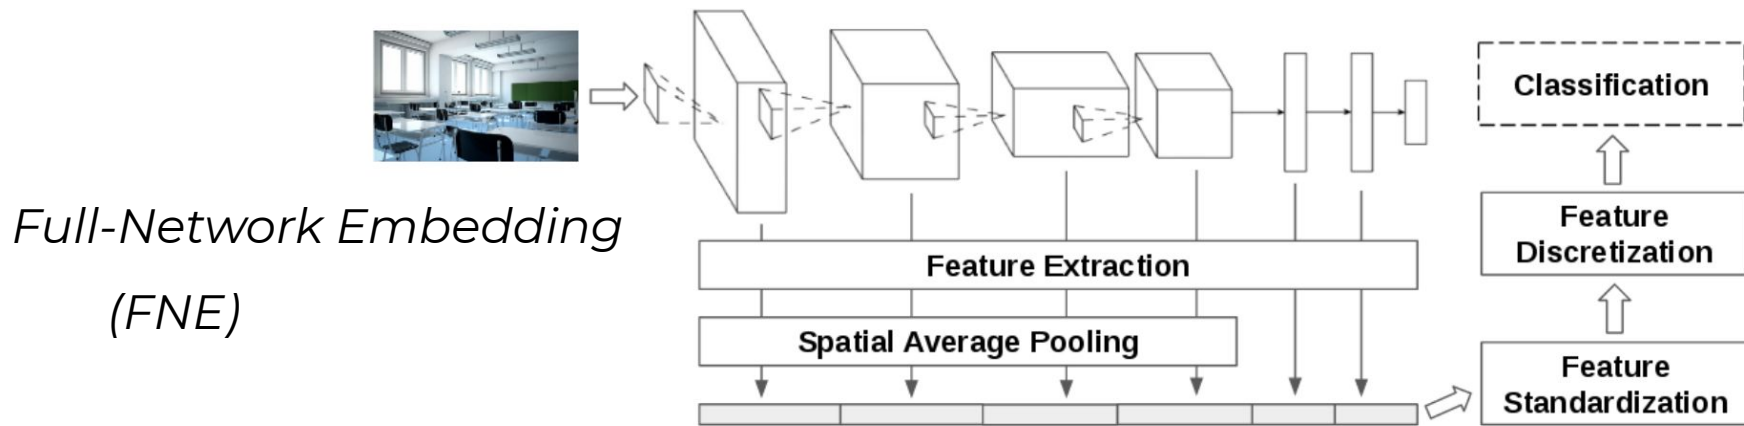
Best for  
multi-layer  
feature  
extraction





# Advanced feature normalization

- Dimensionality of extracted features is an issue (12K in VGG16)
- Removing complexity without losing expressivity
  - Discretizing the space (-1,0,1)



# Feature extraction in action

- High similarity source - target

Network pre-trained on **Places2** for mit67 and on **ImageNet** for the rest.

Dataset	mit67	cub200	flowers102	cats-dogs	sdogs	caltech101	food101	textures	wood
Baseline fc6	80.0	65.8	89.5	89.3	78.0	91.4 $\pm$ 0.6	61.4 $\pm$ 0.2	69.6	70.8 $\pm$ 6.6
Baseline fc7	81.7	63.2	87.0	89.6	79.3	89.7 $\pm$ 0.3	59.1 $\pm$ 0.6	69.0	68.9 $\pm$ 6.8
Full-network	83.6	65.5	93.3	89.2	78.8	91.4 $\pm$ 0.6	67.0 $\pm$ 0.7	73.0	74.1 $\pm$ 6.9
SotA	86.9 [5]	92.3 [10]	97.0 [5]	91.6 [6]	90.3 [5]	93.4 [31]	77.4 [4]	75.5 [17]	- -
ED	✓	✓	✓	✗	✓	✗	✗	✗	-
FT	✓	✓	✓	✓	✓	✓	✓	✗	-

# Feature extraction in action

- High similarity source - target

Network pre-trained on **Places2** for mit67 and on **ImageNet** for the rest.

Dataset	mit67	cub200	flowers102	cats-dogs	sdogs	caltech101	food101	textures	wood
Baseline f c6	80.0	65.8	89.5	89.3	78.0	91.4±0.6	61.4±0.2	69.6	70.8±6.6
Baseline f c7	81.7	63.2	87.0	89.6	79.3	89.7±0.3	59.1±0.6	69.0	68.9 ±6.8
Full-network	83.6	-0.3	93.3	-0.4	-0.5	91.4±0.6	67.0±0.7	73.0	74.1±6.9
SotA	86.9 [5]	92.3 [10]	97.0 [5]	91.6 [6]	90.3 [5]	93.4 [31]	77.4 [4]	75.5 [17]	- -
ED	✓	✓	✓	✗	✓	✗	✗	✗	-
FT	✓	✓	✓	✓	✓	✓	✓	✗	-

+2.9

+4.2

Task similarity makes single layer l2-norm competitive

# Feature extraction in action

- High similarity source - target

Network pre-trained on **Places2** for mit67 and on **ImageNet** for the rest.

Dataset	mit67	cub200	flowers102	cats-dogs	sdogs	caltech101	food101	textures	wood
Baseline fc6	80.0	65.8	89.5	89.3	78.0	91.4±0.6	61.4±0.2	69.6	70.8±6.6
Baseline fc7	81.7	63.2	87.0	89.6	79.3	89.7±0.3	59.1±0.6	69.0	68.9 ±6.8
Full-network	83.6	65.5	93.3	89.2	78.8	91.4±0.6	67.0±0.7	73.0	74.1±6.9
SotA	86.9 [5]	92.3 [10]	97.0 [5]	91.6 [6]	90.3 [5]	93.4 [31]	77.4 [4]	75.5 [17]	- -
ED	✓	✓	✓	✗	✓	✗	✗	✗	-
FT	✓	✓	✓	✓	✓	✓	✓	✗	-

+7.9

**Data (external or not) can make fine tuning worth the COST**

# Feature extraction in action

- Low similarity source - target (*most real-world scenario!*)

Network pre-trained on **ImageNet** for mit67 and on **Places2** for the rest.

Dataset	mit67	cub200	flowers102	cats-dogs	caltech101	textures	wood
Baseline fc7	72.2	23.6	73.3	38.7	72.0	55.8	65.3
Full-network	75.5	35.5	88.7	56.2	80.0	65.1	74.0
	+3.3	+11.9	+15.4	+17.5	+8.0	+9.3	+10.6

**Data (external or not) makes fine tuning worth the COST**

# Key takeaways

- If possible, always use a pre-trained net
  - Don't be a hero
- Consider the gradient of representations
  - From data to task
- Always analyze
  - Source/Target similarity
  - Data availability

# Key takeaways

- Fine tune if possible
  - Freeze from the bottom
  - Fine tune the middle
  - Retrain from random the top
- Feature extraction
  - Must-do baseline (cheap and easy!)
  - Only feasible approach if data volume is short

**Dario Garcia-Gasulla (BSC)**  
*[dario.garcia@bsc.es](mailto:dario.garcia@bsc.es)*

